Explore Data Analytics in Azure

Create a Synapse workspace

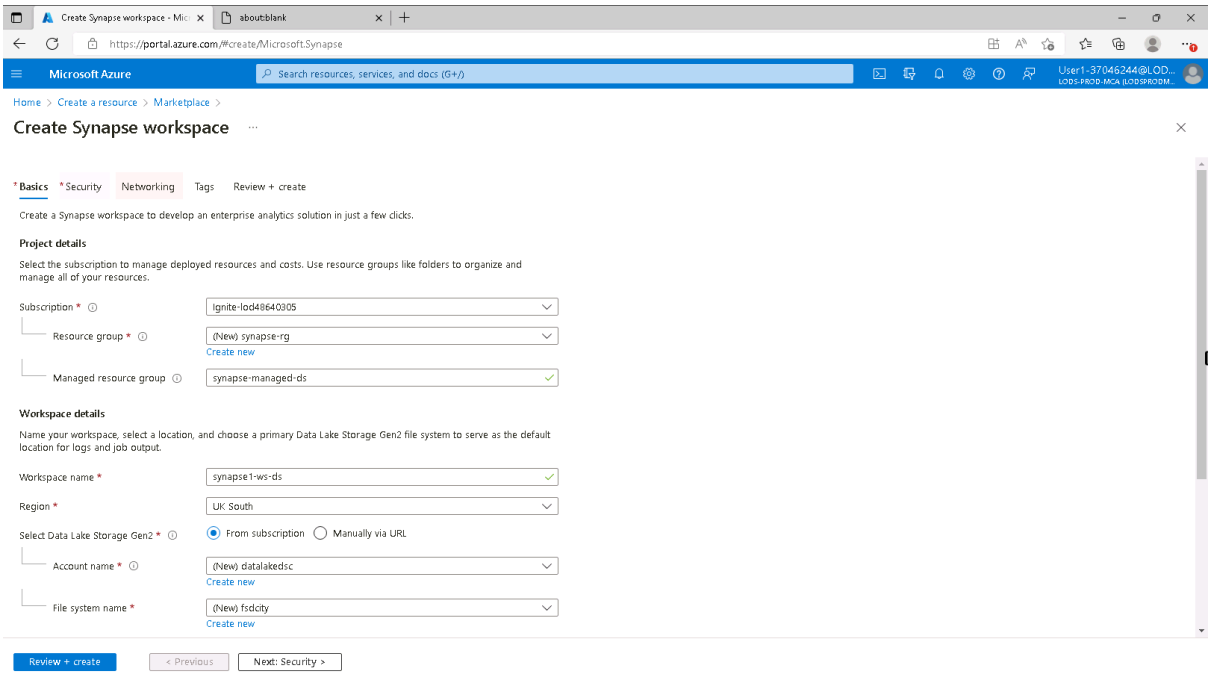

Create a Lakehouse

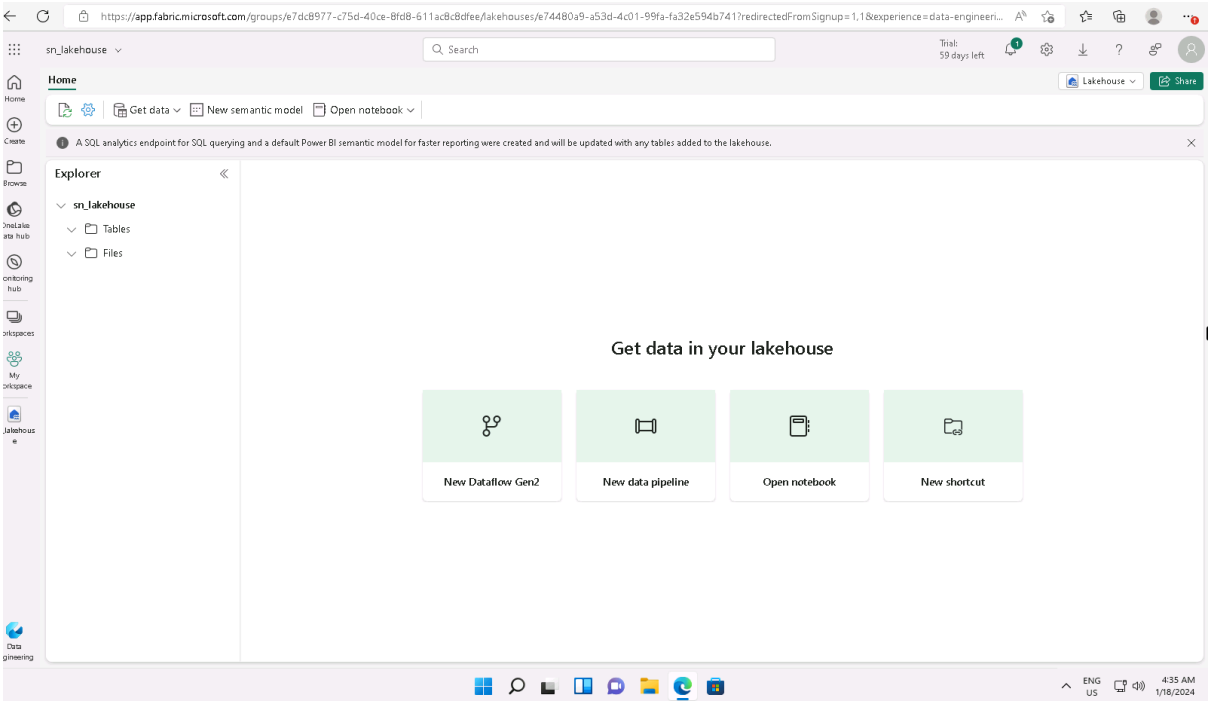

Ingest Data

- A simple way to ingest data is to use a **Copy Data** activity in a pipeline to extract the data from a source and copy it to a file in the lakehouse.

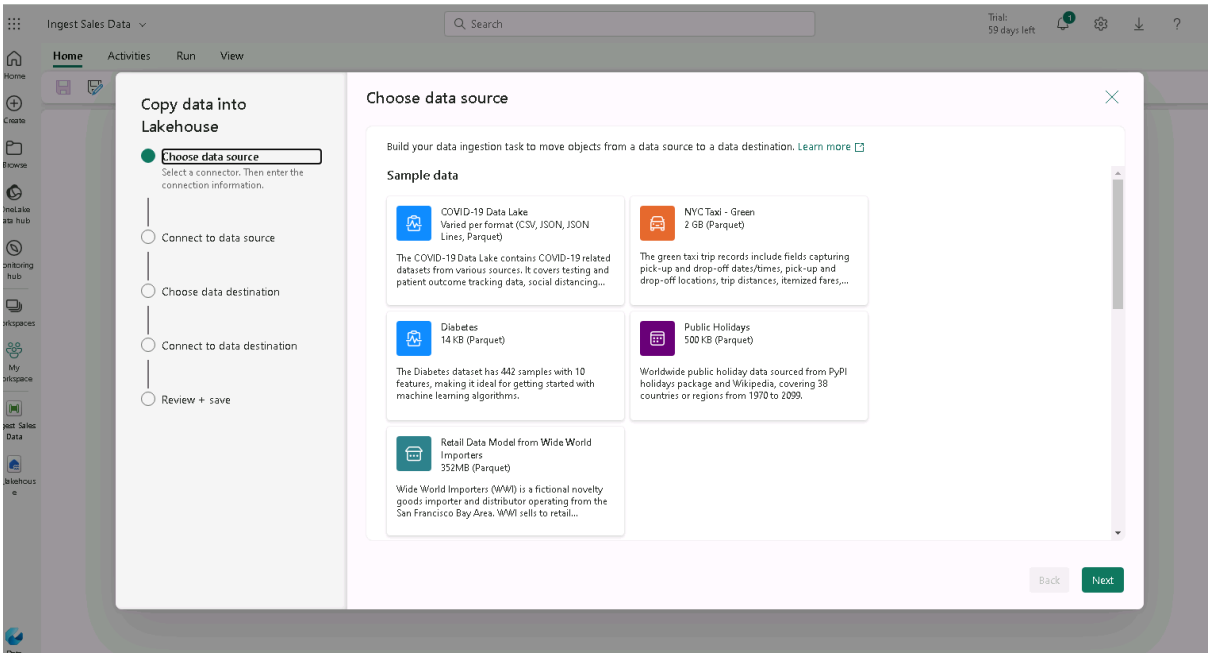

1. Connecting to a data source

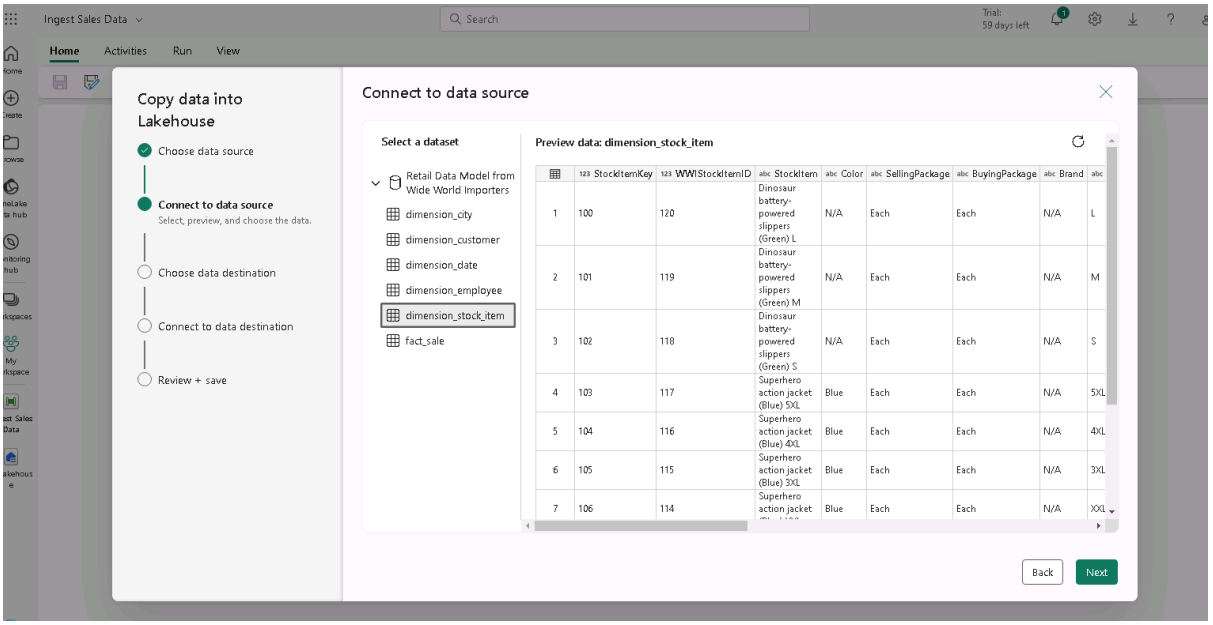

2. Choosing an existing data destination

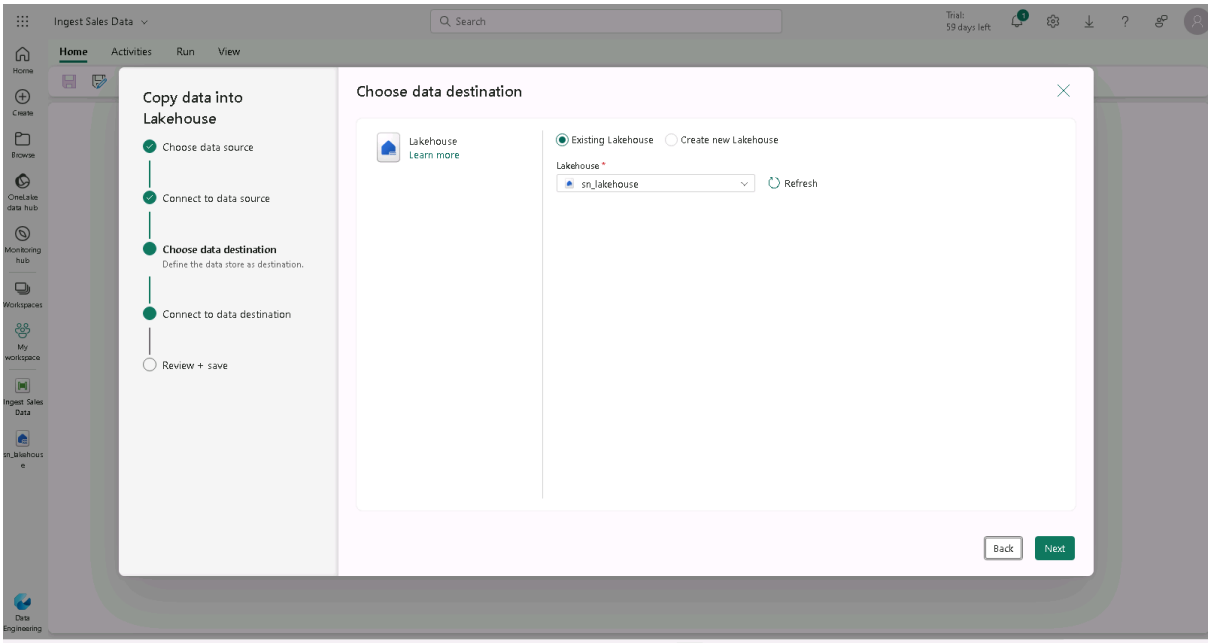
Ingest Sales Data
Search
Trial: 59 days left
Home
Activities
Run
View
Copy data into Lakehouse
Choose data source
Connect to data source
Choose data destination
Define the data store as destination.
Connect to data destination
Review + save
Choose data destination
Lakehouse
Learn more
Existing Lakehouse
Create new Lakehouse
Lakehouse *
sn_lakehouse
Refresh
Back
Next

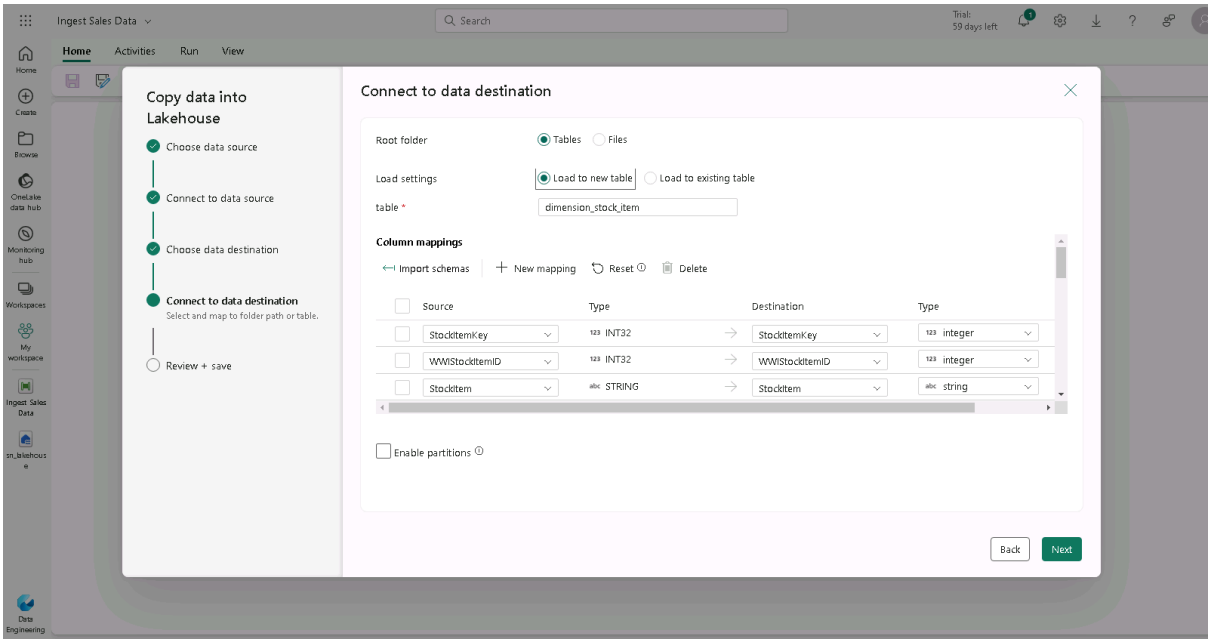
Ingest Sales Data
Search
Trial: 59 days left
Home
Activities
Run
View
Copy data into Lakehouse
Choose data source
Connect to data source
Choose data destination
Connect to data destination
Select and map to folder path or table.
Review + save
Connect to data destination
Root folder
Tables
Files
Load settings
Load to new table
Load to existing table
table *
dimension_stock_item
Column mappings
Import schemas
New mapping
Reset
Delete
Source
Type
Destination
Type
StockItemKey
INT32
StockItemKey
integer
WWIStockItemID
INT32
WWIStockItemID
integer
StockItem
STRING
StockItem
string
Enable partitions
Back
Next

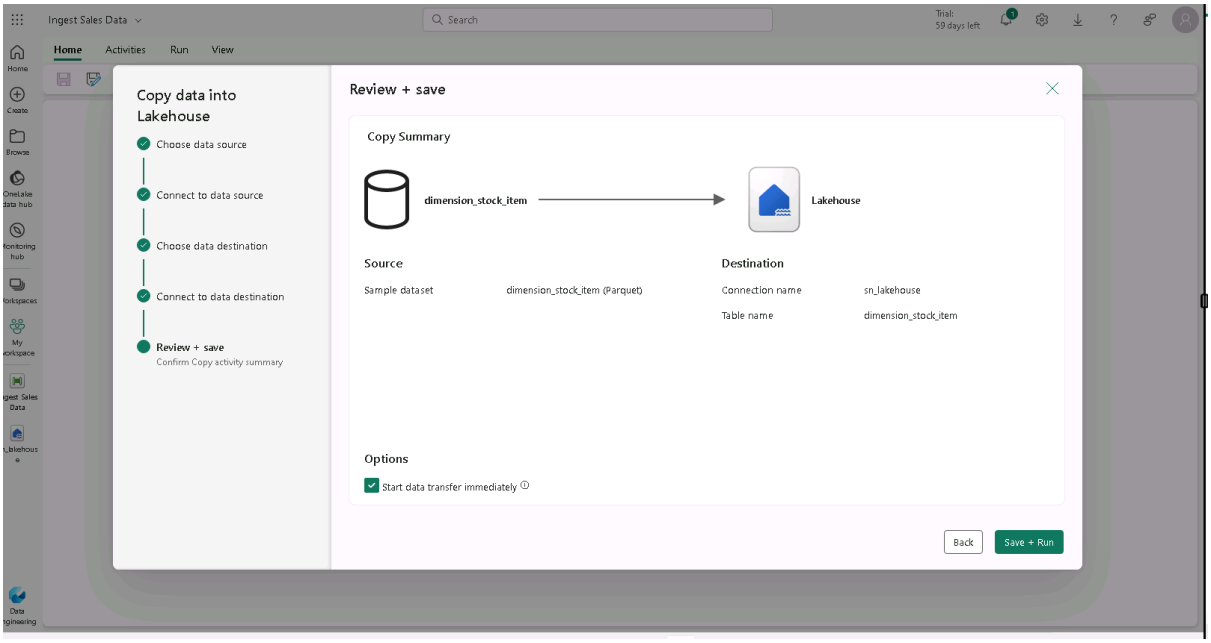

A new pipeline containing a **Copy Data** activity is created, as shown here

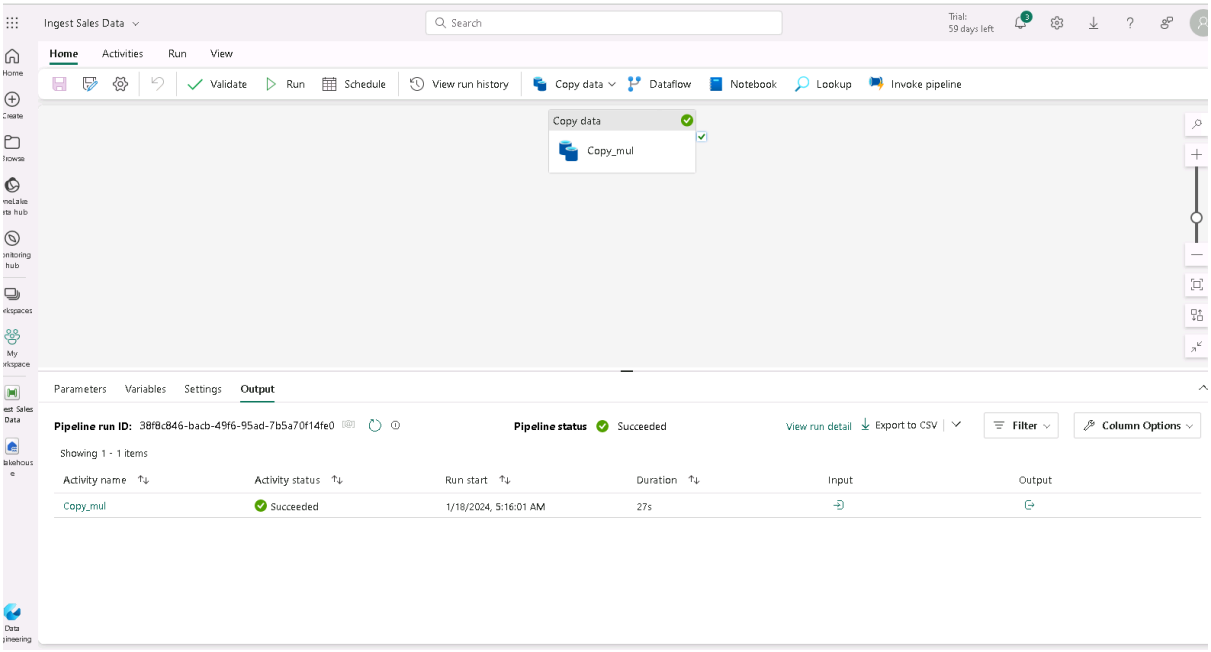

Results:

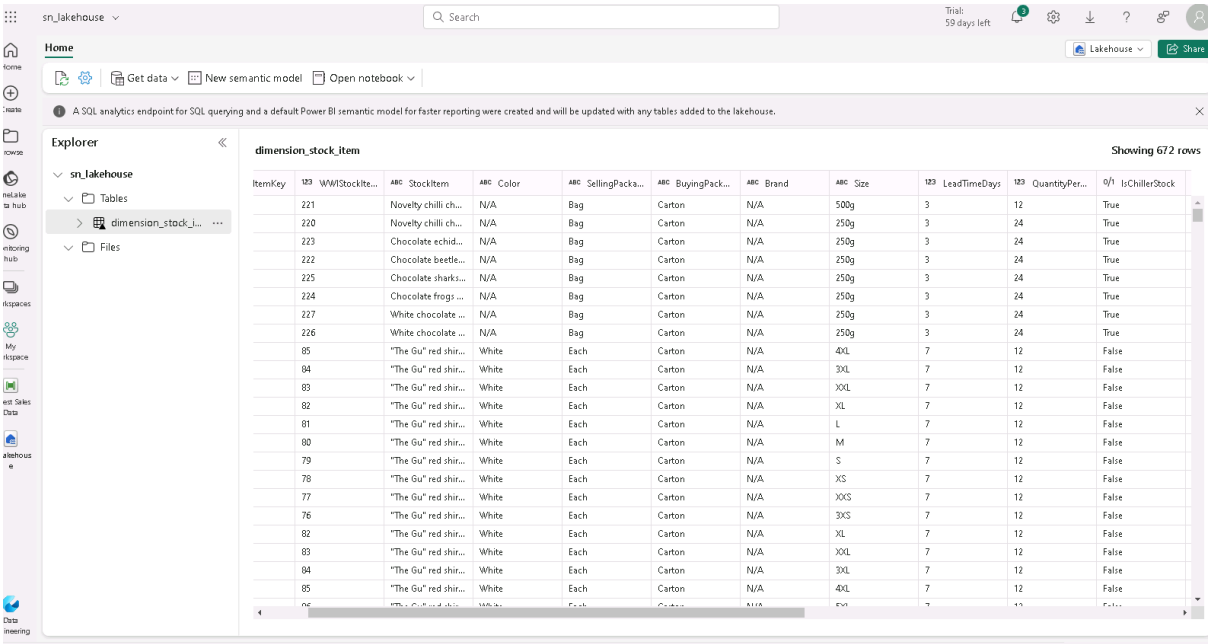

Query data in a lakehouse:

1. Select SQL analytical endpoint

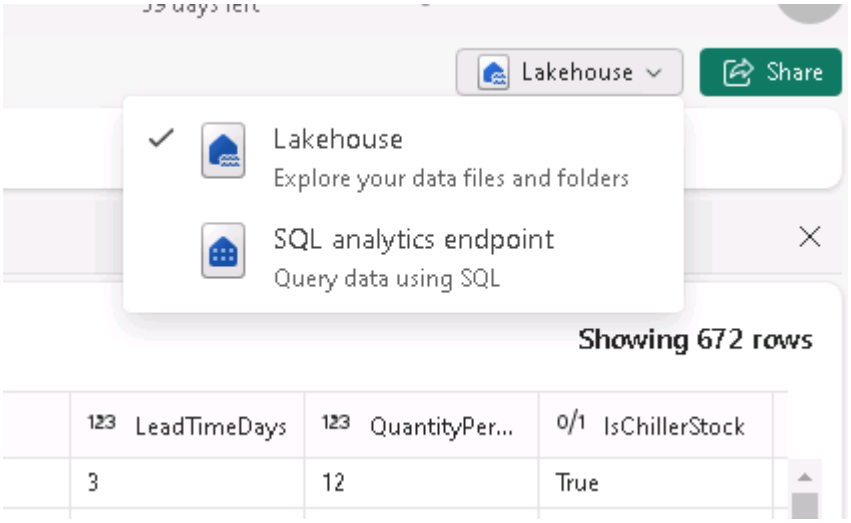

2. Select new sql query and entered the sql code into the query editor. Ran the query and it revealed that there are two brand values - N/A and Northwide

and shows the number of products in each

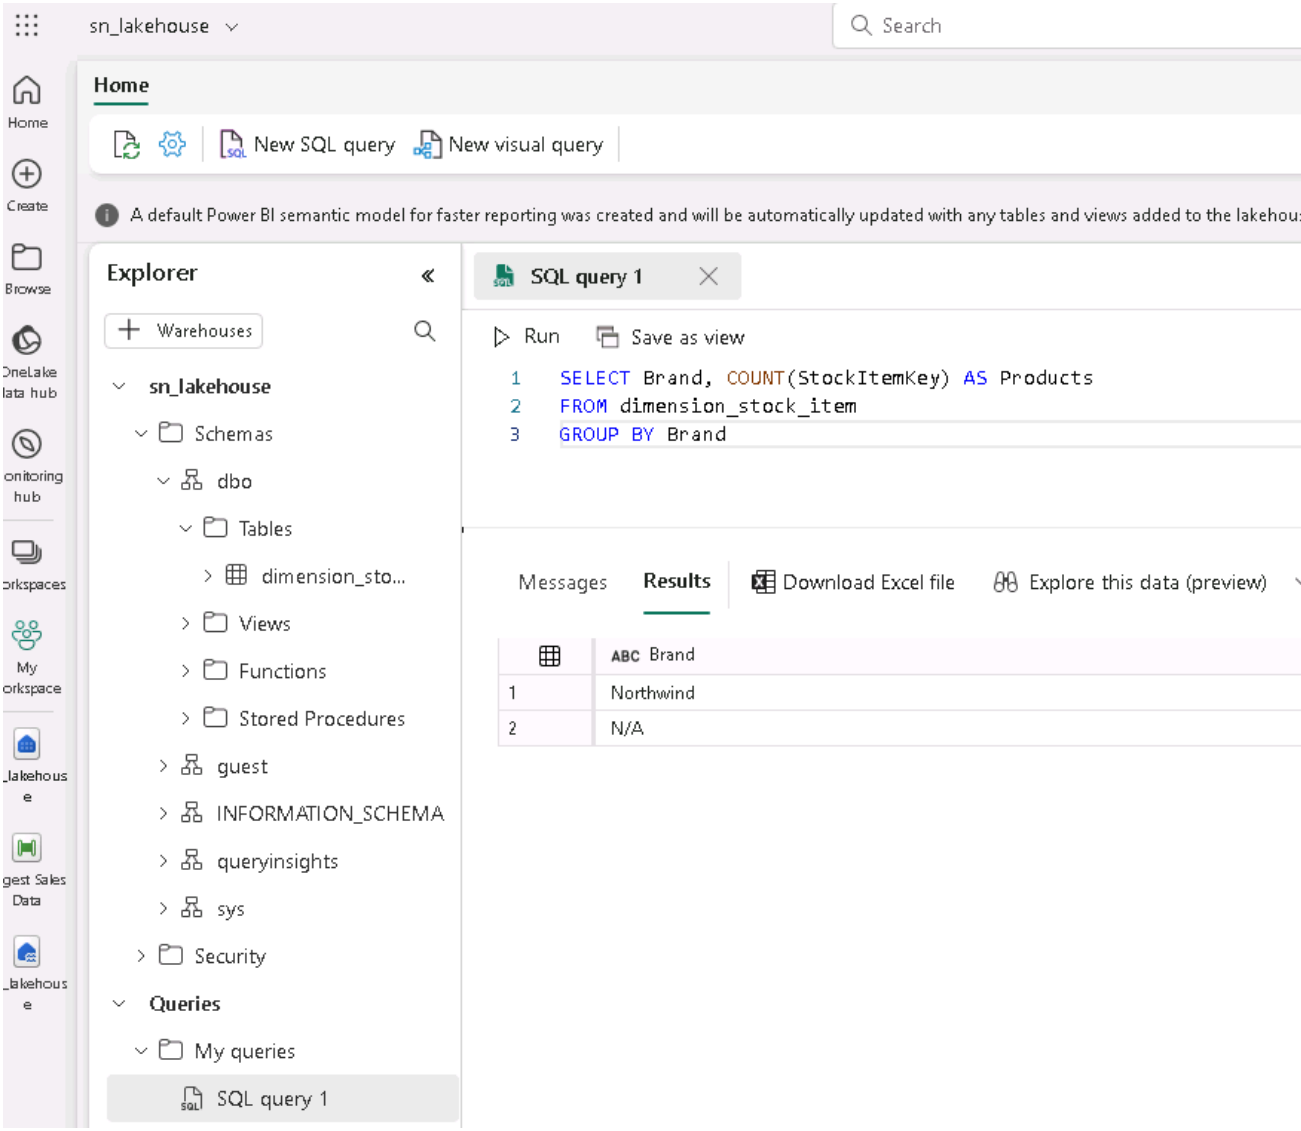

Visualise you data in a lakehouse:

1. Select model tab. This will let you see the data for the tables in the lakehouse.

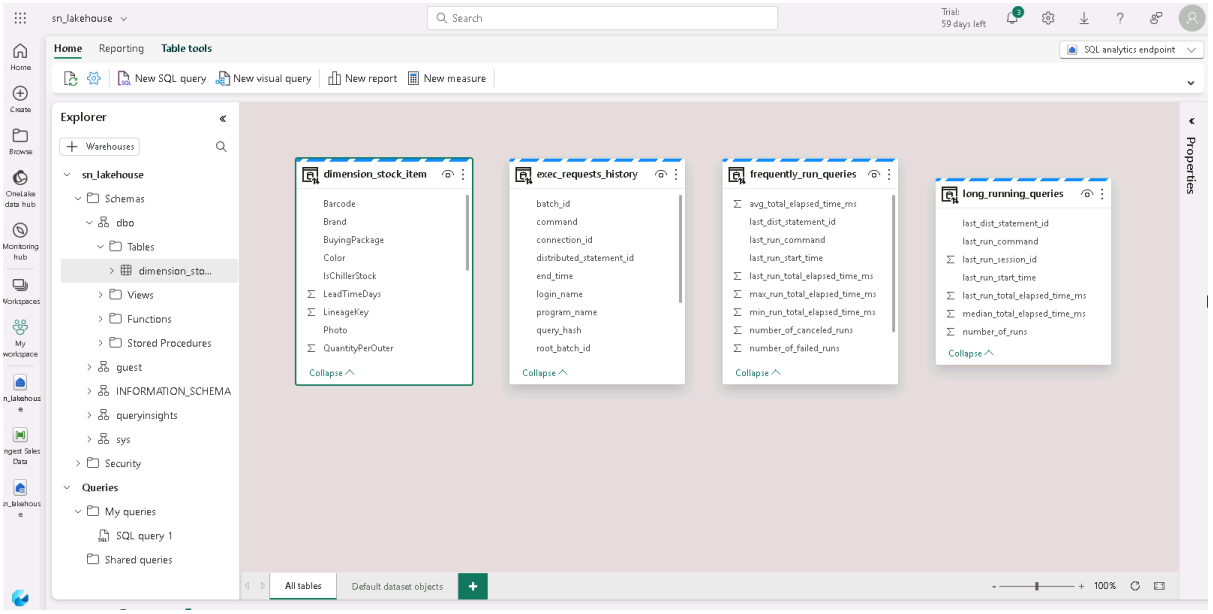

2. Selected stacked barchart and input brand and stockitemkey fields. Then the aggregation in the y-axis and x-axis was changed to count. Next, added colours and resized the chart to make it more visualsing.

3. 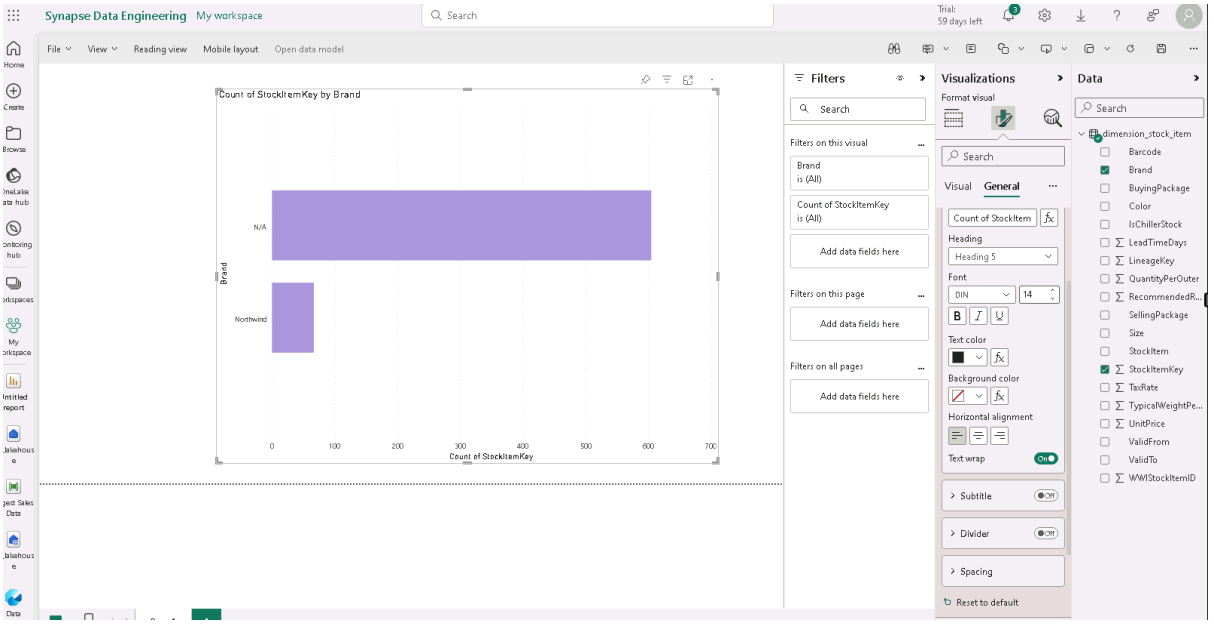

4. Saved it as brand quantity report

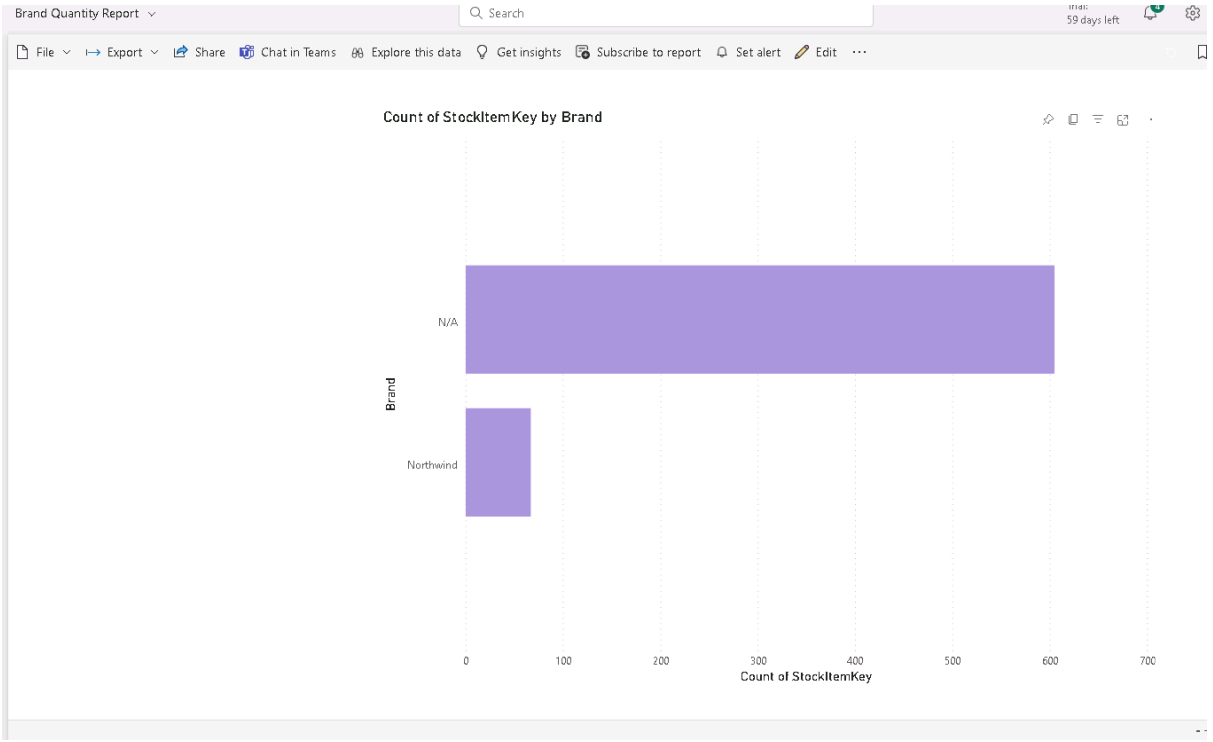

Create a KQL Database

1. Switched to Real Time Analytics. The real time analytics home page includes tiles to create commonly used assets for real-time data analyis. Named the database ive created ' my_kql_database'.

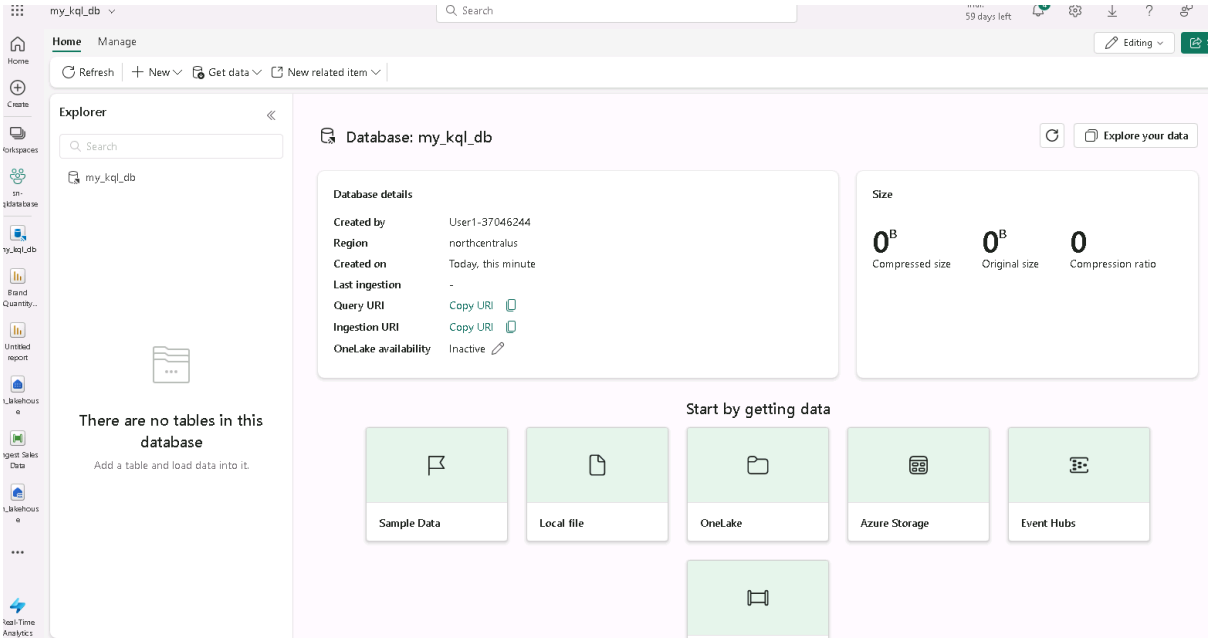

Create an eventstream

1. Eventstreams provide a scalable and flexible way to ingest real-time data from a streaming source.

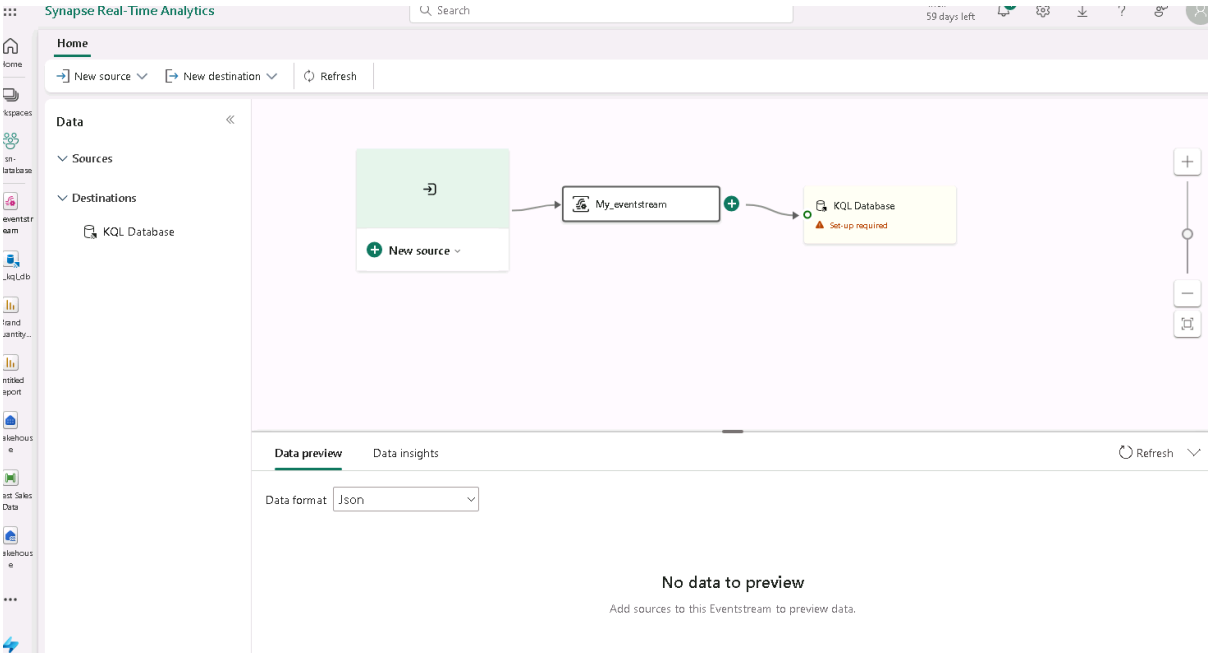

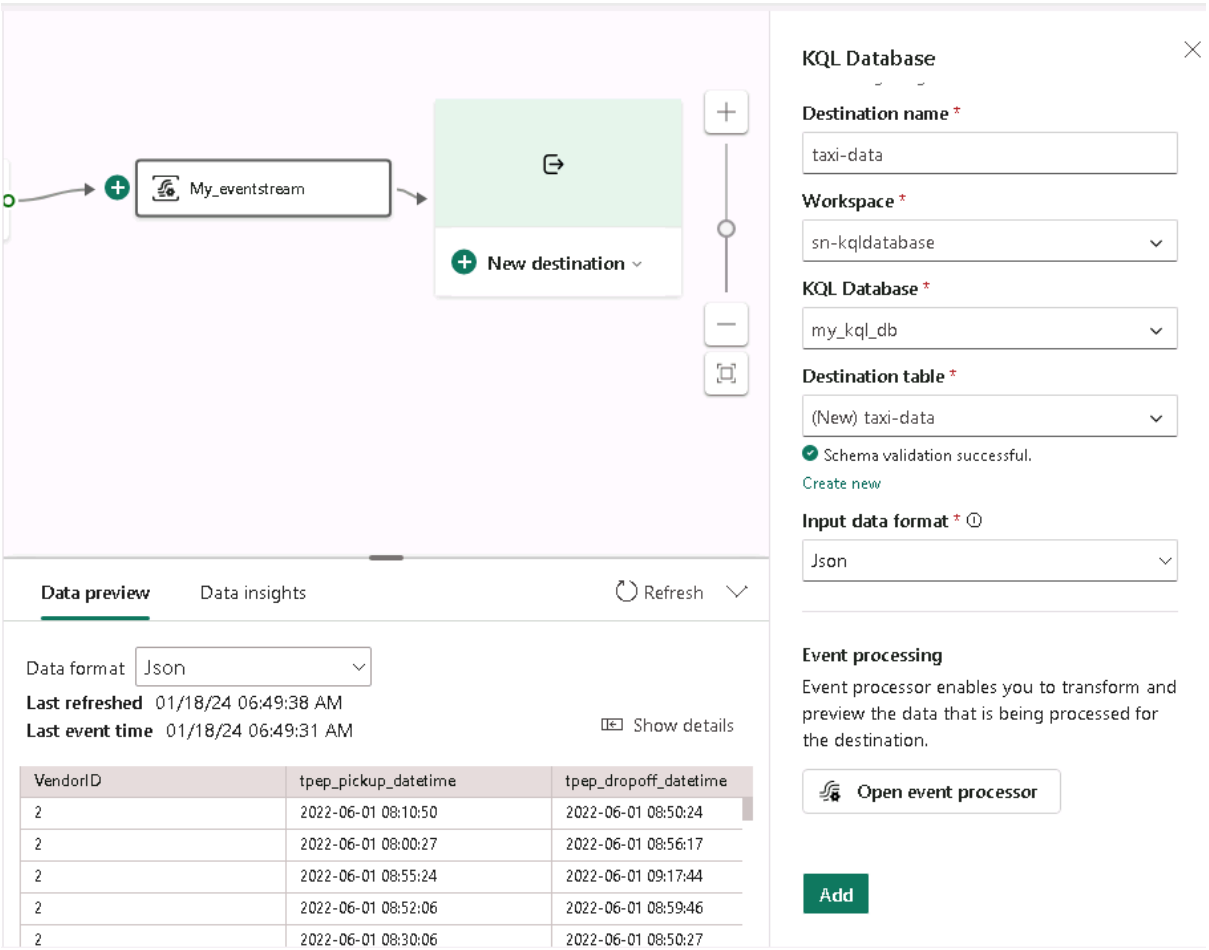
My_eventstream
New destination
KQL Database
Destination name *
taxi-data
Workspace *
sn-kqldatabase
KQL Database *
my_kql_db
Destination table *
(New) taxi-data
Schema validation successful.
Create new
Input data format *
Json
Data preview
Data insights
Refresh
Data format
Json
Last refreshed 01/18/24 06:49:38 AM
Last event time 01/18/24 06:49:31 AM
Show details
VendorID
tpep_pickup_datetime
tpep_dropoff_datetime
2
2022-06-01 08:10:50
2022-06-01 08:50:24
2
2022-06-01 08:00:27
2022-06-01 08:56:17
2
2022-06-01 08:55:24
2022-06-01 09:17:44
2
2022-06-01 08:52:06
2022-06-01 08:59:46
2
2022-06-01 08:30:06
2022-06-01 08:50:27
Event processing
Event processor enables you to transform and preview the data that is being processed for the destination.
Open event processor
Add

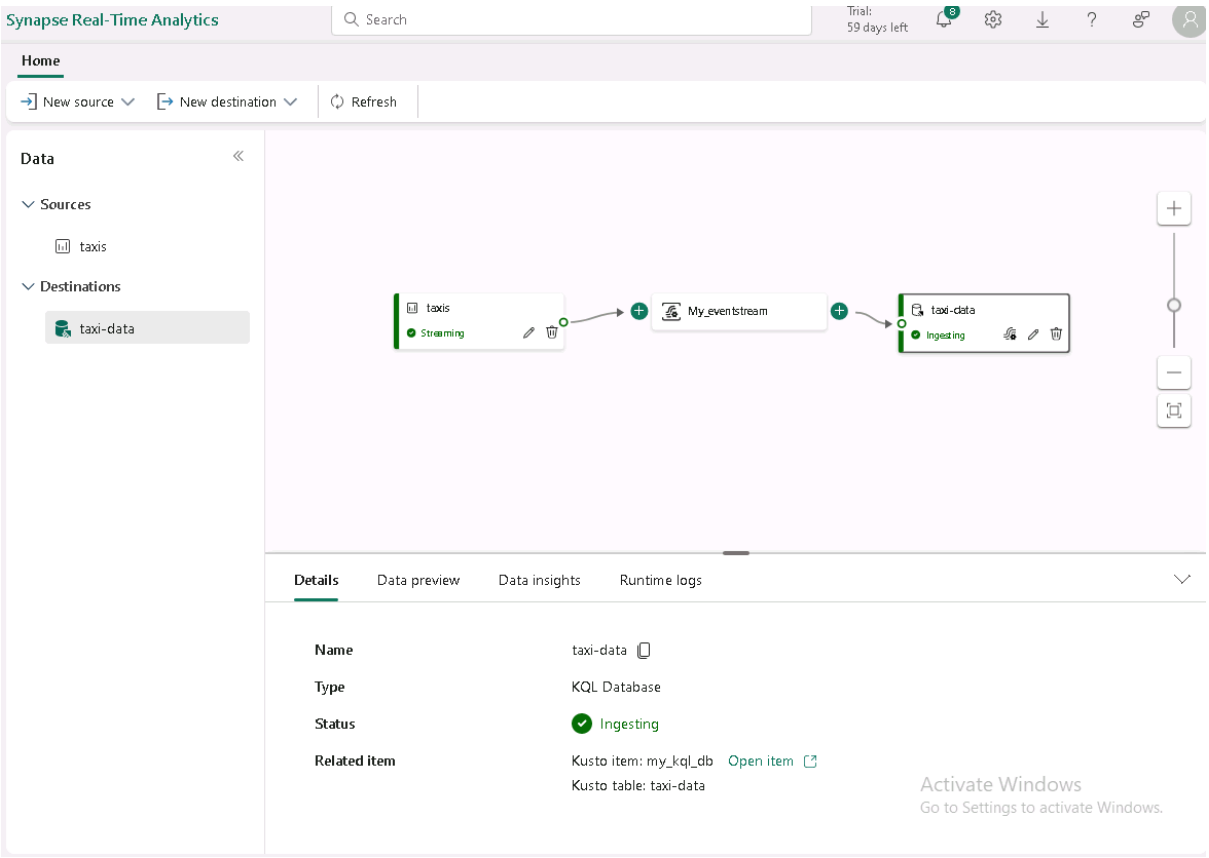
Synapse Real-Time Analytics
Search
Trial:
59 days left
Home
New source
New destination
Refresh
Data
Sources
taxis
Destinations
taxi-data
taxis
Streaming
My_eventstream
taxi-data
Ingesting
Details
Data preview
Data insights
Runtime logs
Name
taxi-data
Type
KQL Database
Status
Ingesting
Related item
Kusto item: my_kql_db
Open item
Kusto table: taxi-data
Activate Windows
Go to Settings to activate Windows.

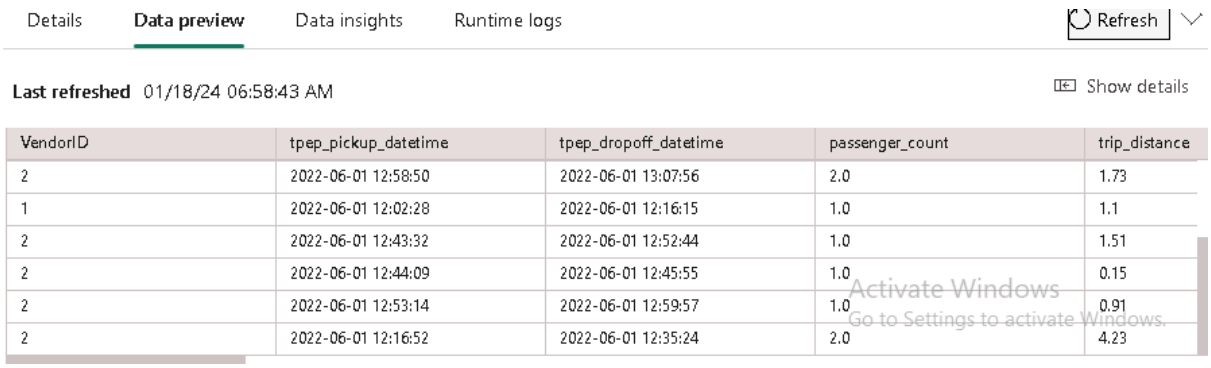

Query real time data in in a KQL data base:

1. select **Query table > Records ingested in the last 24 hours**.

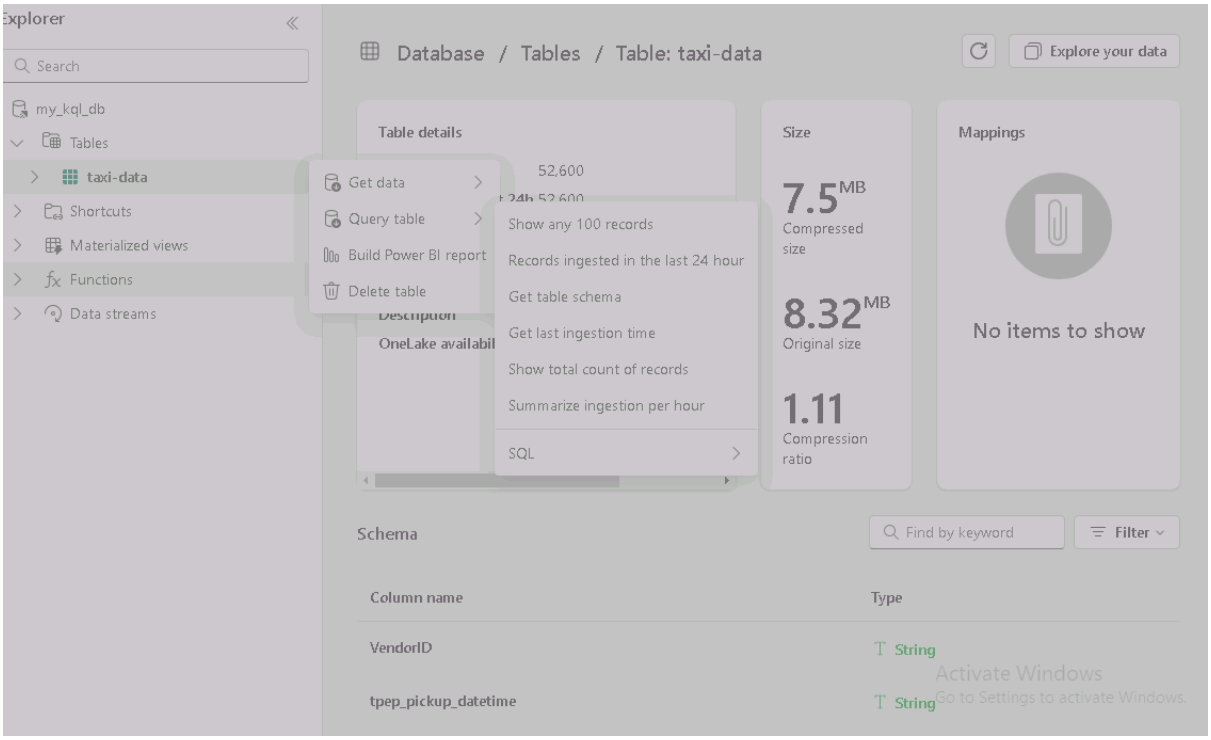

2. This KQL query shows all taxi records ingested from the streaming source in the last 24 hours.

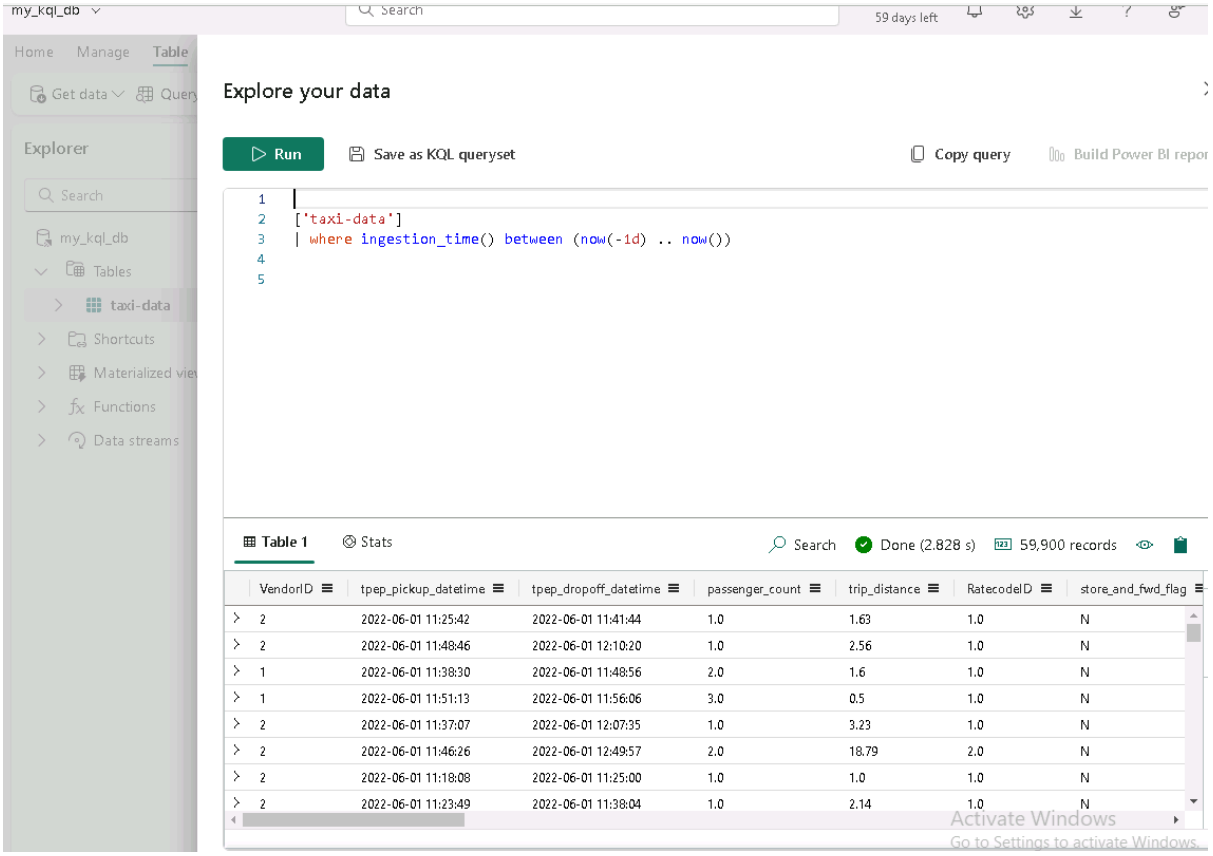

3. This KQL query code shows the number of taxi pickups for each hour.

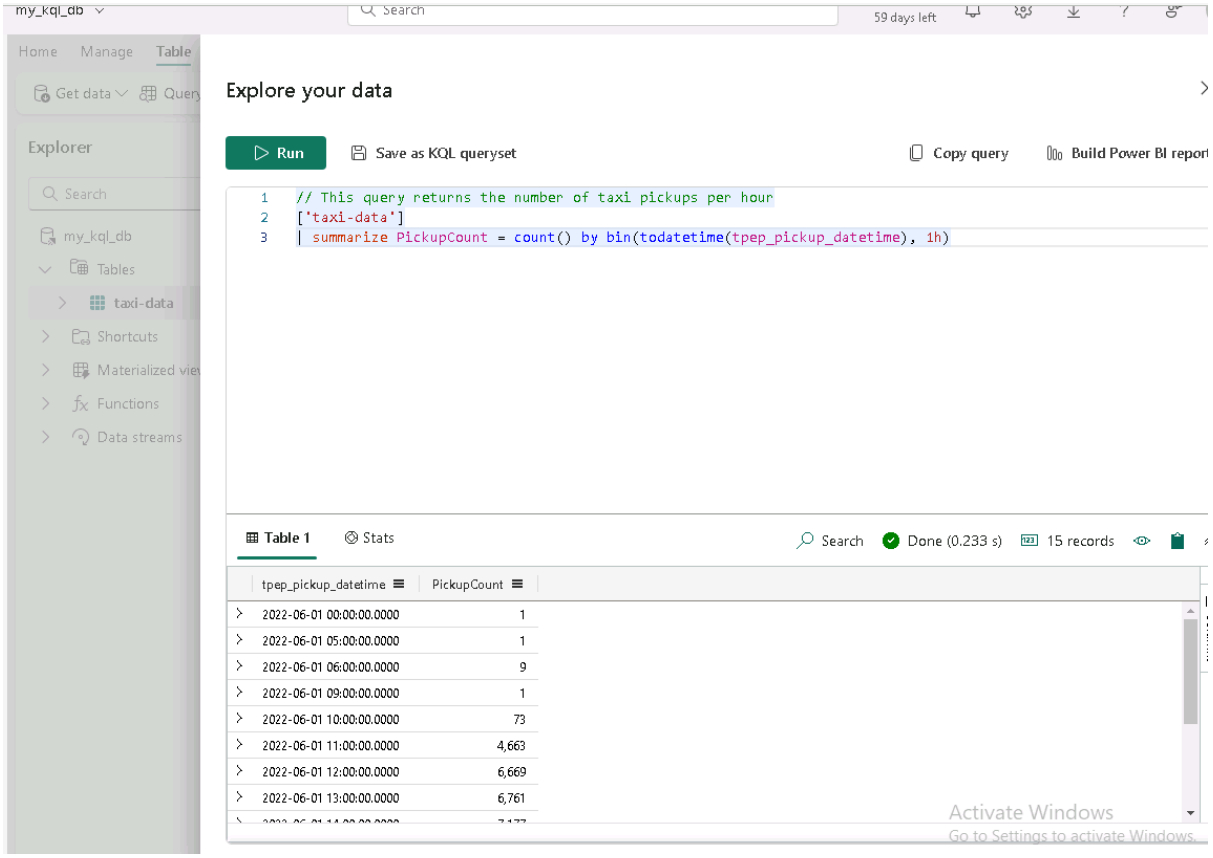
my_kql_db
Search
59 days left
Home
Manage
Table
Get data
Explorer
Search
my_kql_db
Tables
taxi-data
Shortcuts
Functions
Data streams
Explore your data
Run
Save as KQL queryset
Copy query
Build Power BI report
// This query returns the number of taxi pickups per hour
['taxi-data']
| summarize PickupCount = count() by bin(todatetime(tpep_pickup_datetime), 1h)
Table 1
Stats
Search
Done (0.233 s)
15 records
tpep_pickup_datetime
PickupCount
2022-06-01 00:00:00.0000 1
2022-06-01 05:00:00.0000 1
2022-06-01 06:00:00.0000 9
2022-06-01 09:00:00.0000 1
2022-06-01 10:00:00.0000 73
2022-06-01 11:00:00.0000 4,663
2022-06-01 12:00:00.0000 6,669
2022-06-01 13:00:00.0000 6,761
Activate Windows
Go to Settings to activate Windows.

Extension: Fabric Exercise

Use delta tables in Apache Spark

Module: Work with Delta Lake tables in Microsoft Fabric

1. Create a lakehouse and upload data
    a. Named new lakehouse
    b. Download the data file for this exercise from `https://github.com/MicrosoftLearning/dp-data/raw/main/products.csv`, saving it as products.csv on your local compute

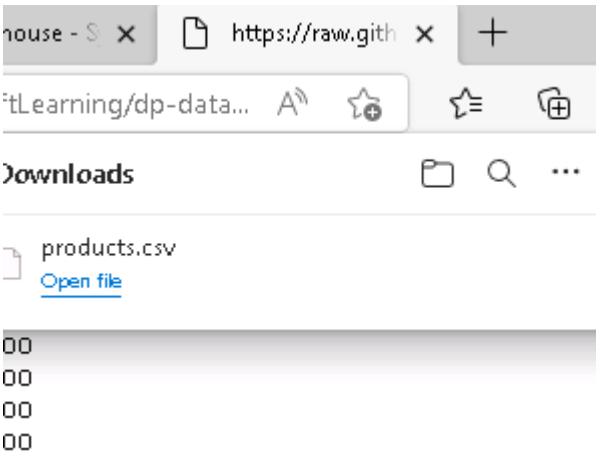

    c. select New subfolder and create a folder named products.
    d. Upload saved file from downloads.

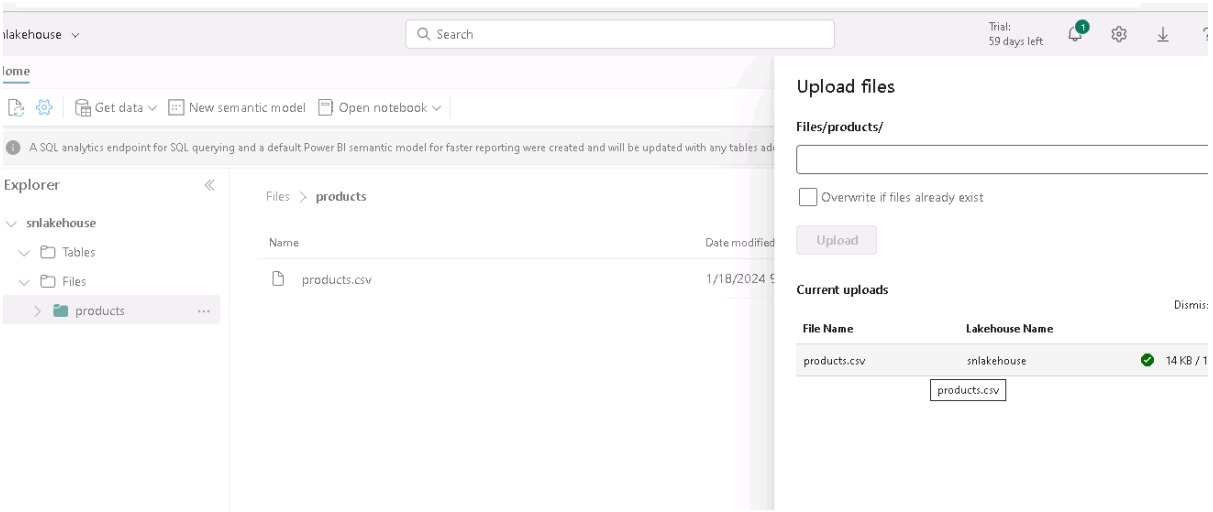

2. Explore data in a dataframe

    a. On the Home page while viewing the contents of the products folder in your datalake, in the Open notebook menu, select New notebook.
    After a few seconds, a new notebook containing a single *cell* will open.

Notebooks are made up of one or more cells that can contain *code* or *markdown* (formatted text).

b. Select the existing cell in the notebook, which contains some simple code, and then use its 🗑 (*Delete*) icon at its top-right to remove it - you will not need this code.

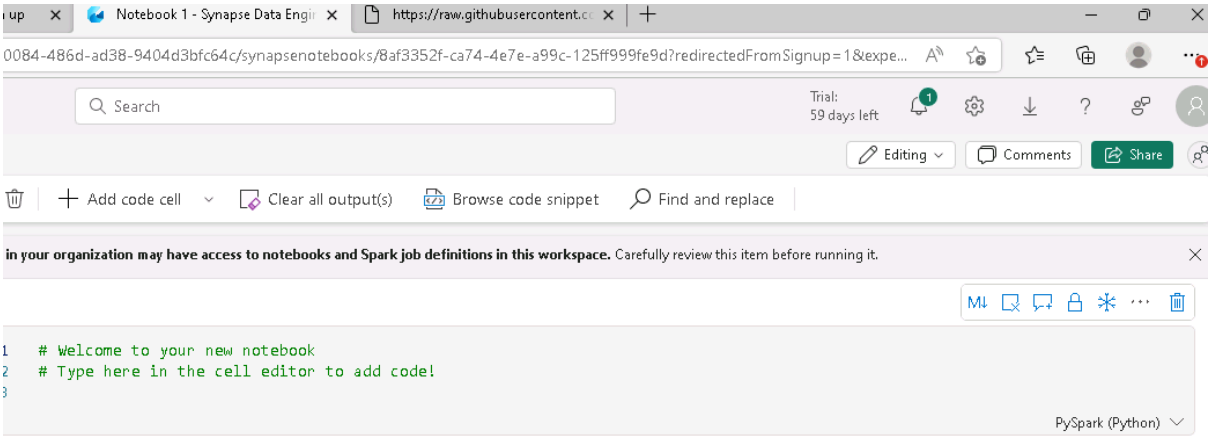

c. In the Lakehouse explorer pane on the left, expand Files and select products to reveal a new pane showing the products.csv file you uploaded previously:

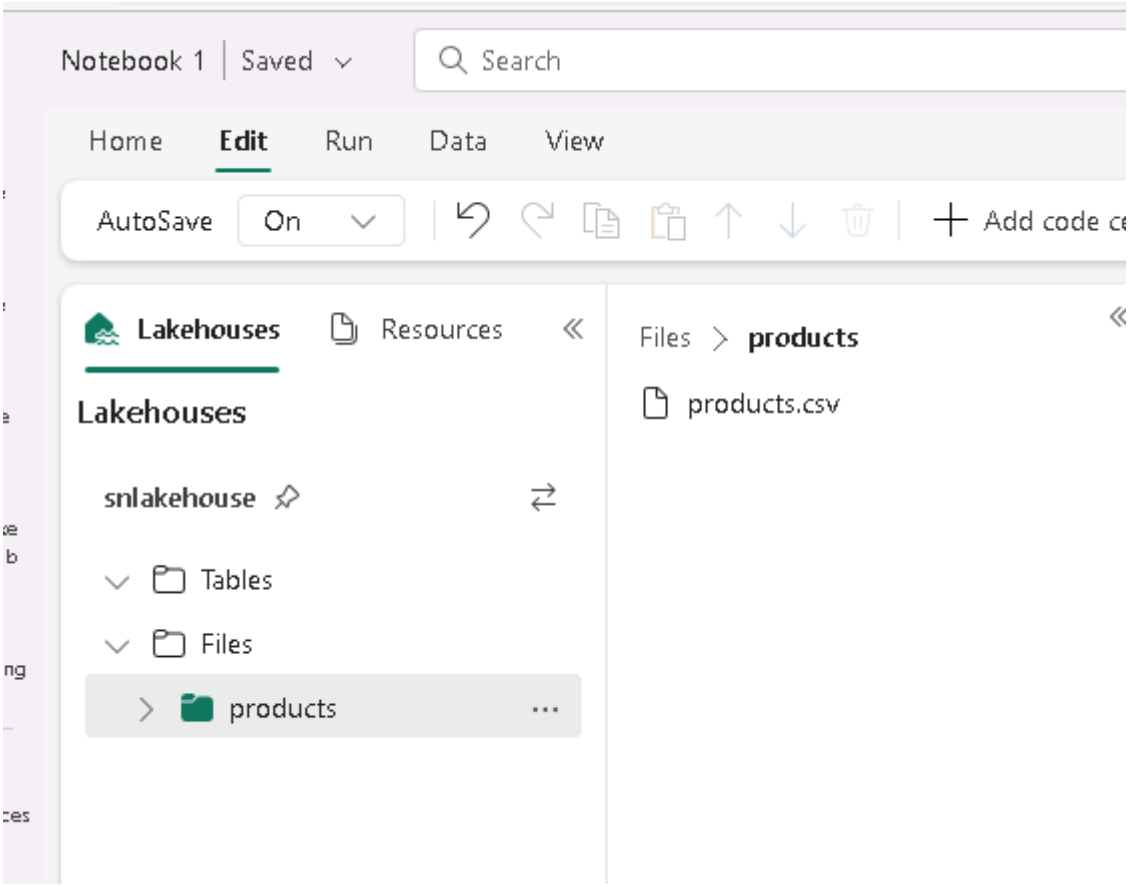

d. In the … menu for products.csv, select Load data > Spark. A new code cell containing the following code should be added to the notebook: Run it.

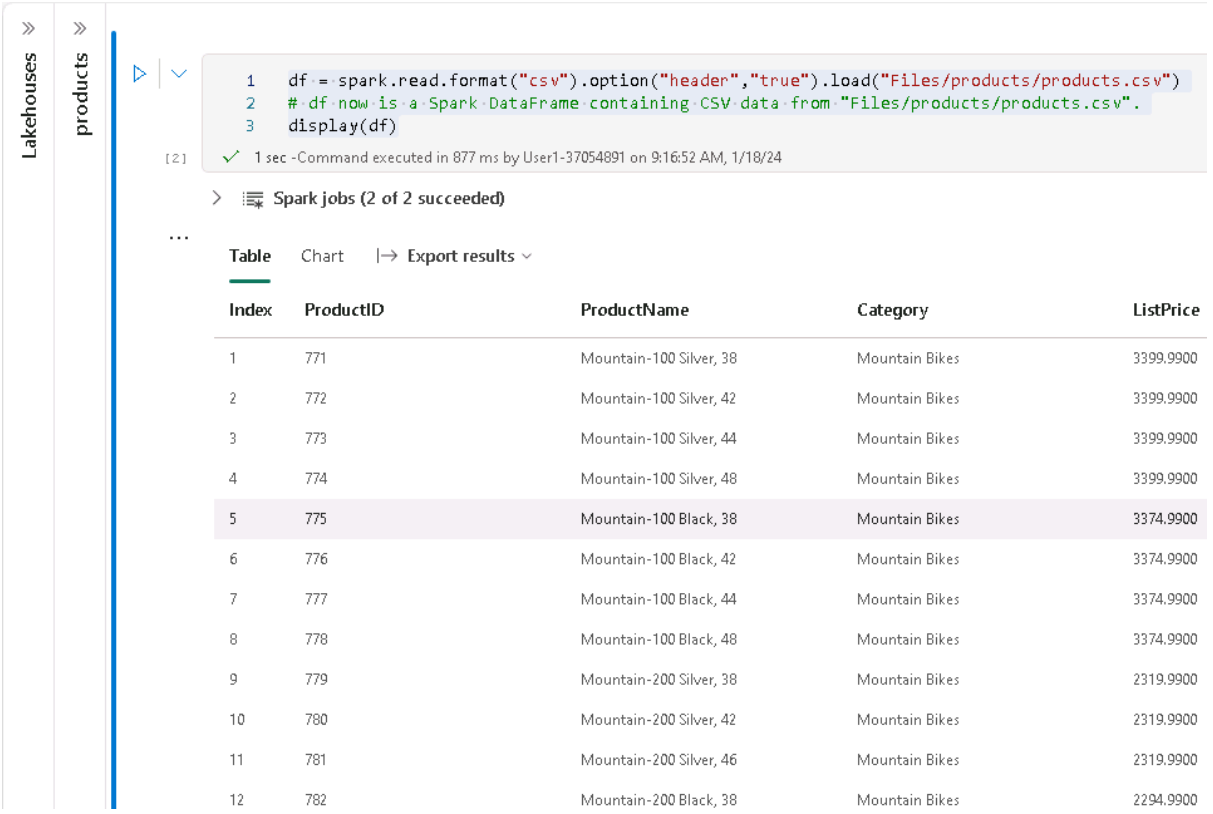

Index	ProductID	ProductName	Category	ListPrice
1	771	Mountain-100 Silver, 38	Mountain Bikes	3399.9900
2	772	Mountain-100 Silver, 42	Mountain Bikes	3399.9900
3	773	Mountain-100 Silver, 44	Mountain Bikes	3399.9900
4	774	Mountain-100 Silver, 48	Mountain Bikes	3399.9900
5	775	Mountain-100 Black, 38	Mountain Bikes	3374.9900
6	776	Mountain-100 Black, 42	Mountain Bikes	3374.9900
7	777	Mountain-100 Black, 44	Mountain Bikes	3374.9900
8	778	Mountain-100 Black, 48	Mountain Bikes	3374.9900
9	779	Mountain-200 Silver, 38	Mountain Bikes	2319.9900
10	780	Mountain-200 Silver, 42	Mountain Bikes	2319.9900
11	781	Mountain-200 Silver, 46	Mountain Bikes	2319.9900
12	782	Mountain-200 Black, 38	Mountain Bikes	2294.9900

Create delta tables

Create a *managed* table; *Managed* tables are tables for which both the schema metadata and the data files are managed by Fabric. The data files for the table are created in the Tables folder.

a. Under the results returned by the first code cell, use the + Code button to add a new code cell if one doesn't already exist. Then enter the following code in the new cell and run it:

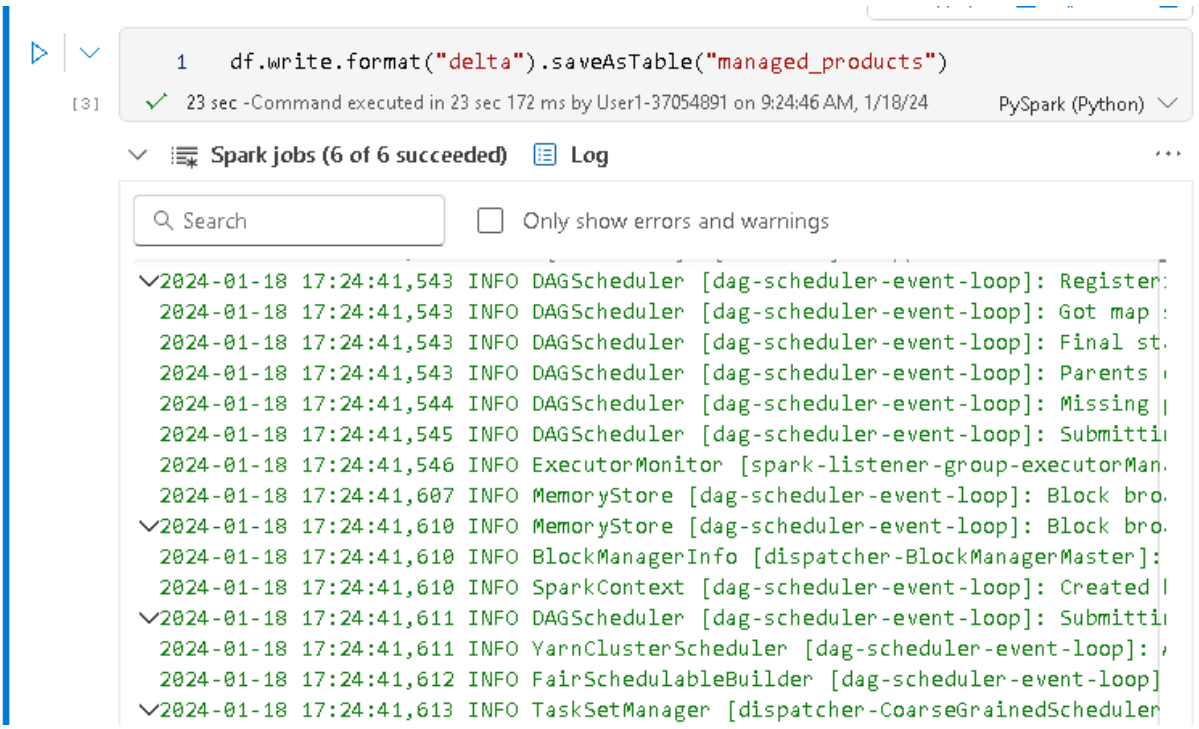

b. In the Lakehouse explorer pane, in the … menu for the Tables folder, select Refresh. Then expand the Tables node and verify that the managed_products table has been created.

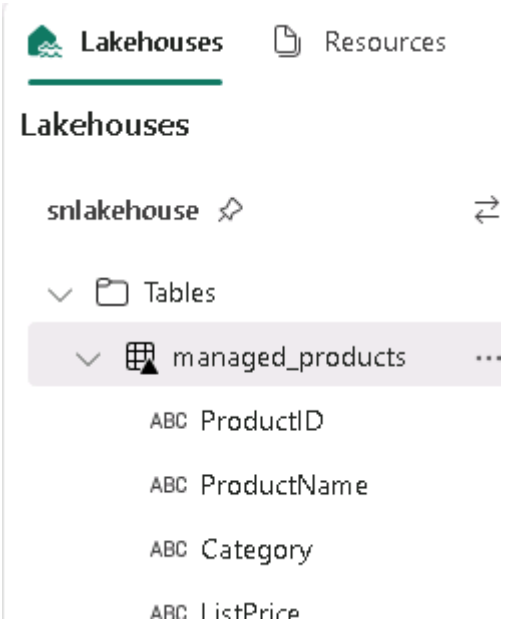

Create an *external* table

You can also create *external* tables for which the schema metadata is defined in the metastore for the lakehouse, but the data files are stored in an external location.

a. Add another new code cell, and add the following code to it:

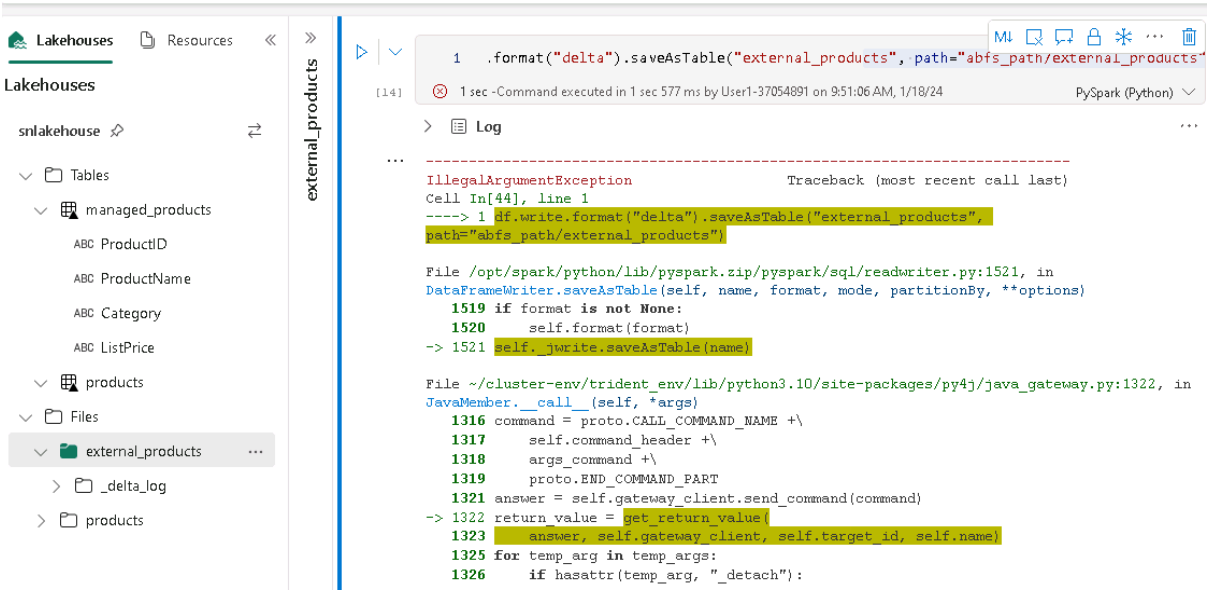

b. In the code you entered into the code cell, replace abfs_path with the path you copied to the clipboard so that the code saves the dataframe as an external table with data files in a folder named external_products in your Files folder location

Create a medallion architecture in a Microsoft Fabric lakehouse

Module: Organize a Fabric lakehouse using medallion architecture design

- Microsoft Fabric is an end-to-end analytics platform that provides a single, integrated environment for data professionals and the business to collaborate on data projects. Fabric provides a set of integrated services that enable you to ingest, store, process, and analyse data in a single environment.

1. Create a workspace: Select synapse data engineering, create a new workspace, navigate to the workspace setting and enable ‘data modelling editing’.

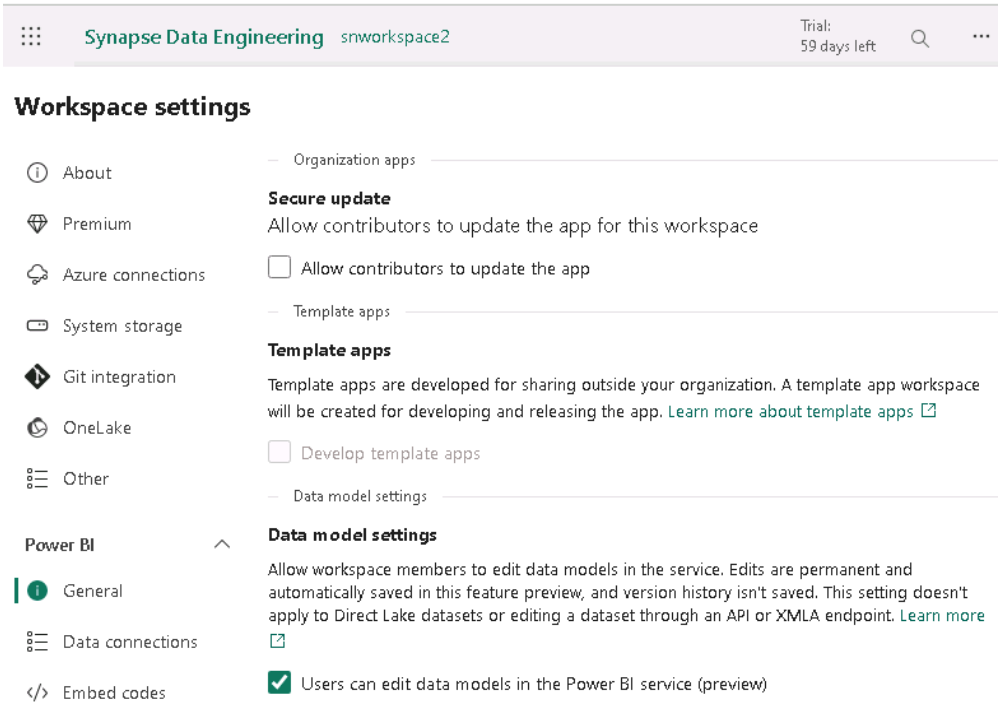

2. Create a lakehouse and upload data to bronze layer

a. created a new lakehouse named sales, extracted folder from github containing three sales data.

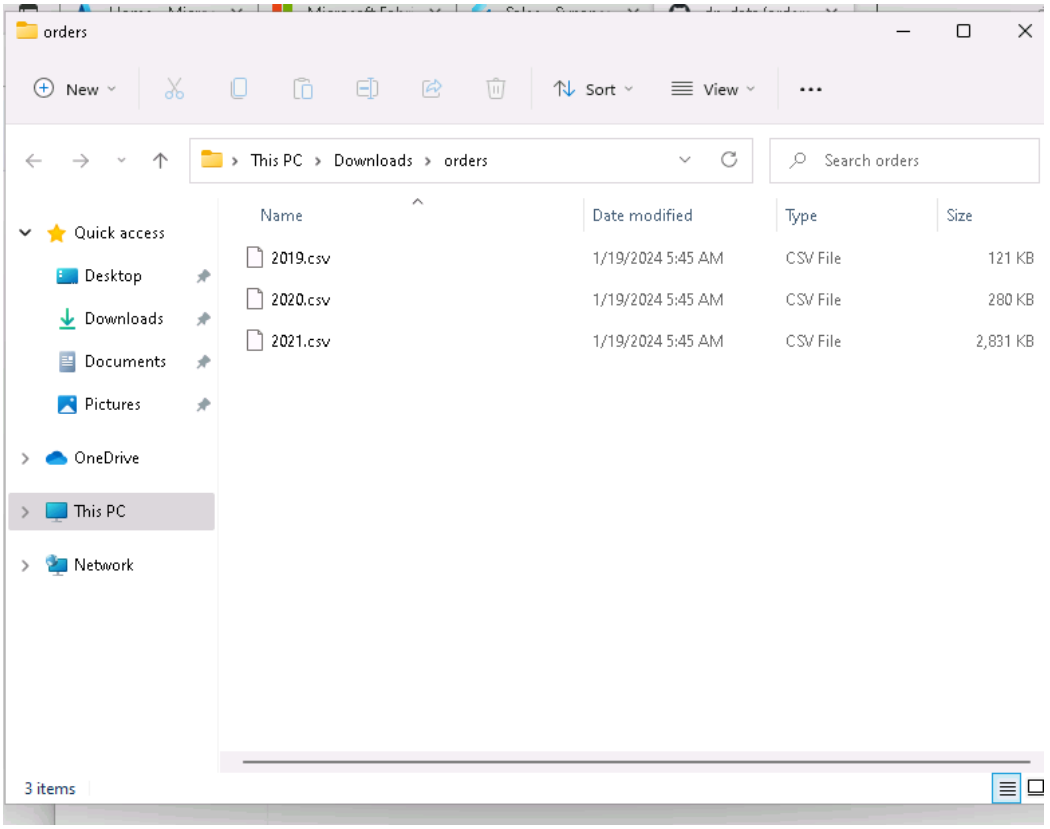

b. Created a new subfolder names bronze under files and upload the extracted files.

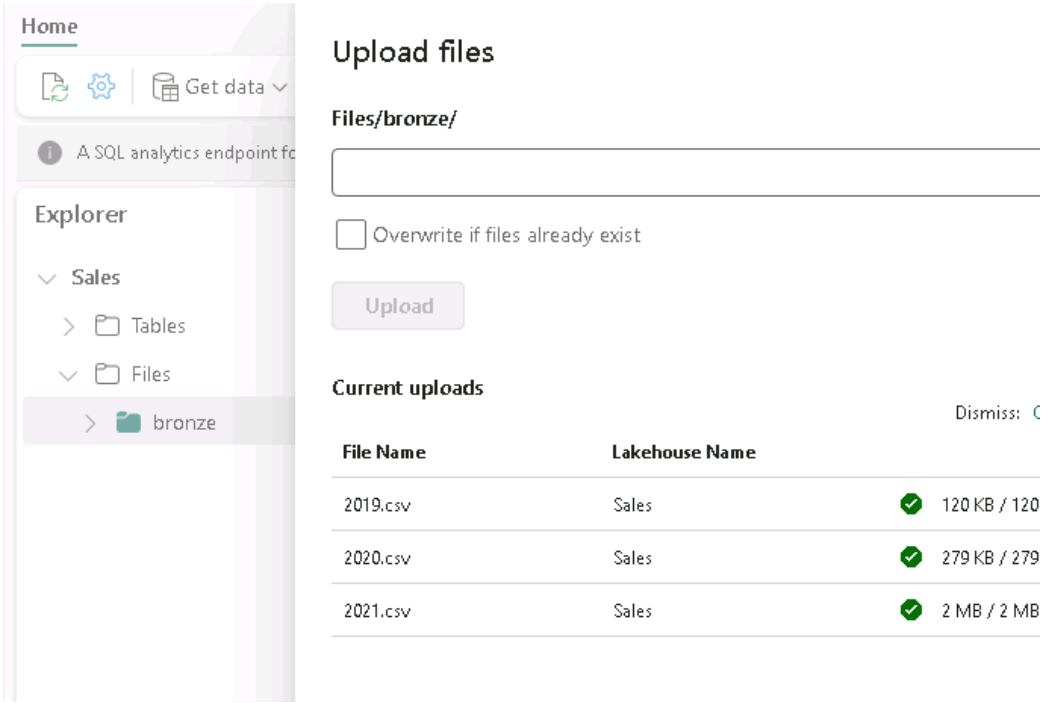

3. Transform data and load to silver delta table

a. Created a new notebook named transform data for silver and inserted a new sparks code into the code cell and ran it. The code loaded the data from the CSV files in the bronze folder into a Spark dataframe, and then displayed the first 10 rows of the dataframe.

The code:

```
from pyspark.sql.types import *

# Create the schema for the table
orderSchema = StructType([
    StructField("SalesOrderNumber", StringType()),
    StructField("SalesOrderLineNumber", IntegerType()),
    StructField("OrderDate", DateType()),
    StructField("CustomerName", StringType()),
    StructField("Email", StringType()),
    StructField("Item", StringType()),
    StructField("Quantity", IntegerType()),
    StructField("UnitPrice", FloatType()),
    StructField("Tax", FloatType())
    ])

# Import all files from bronze folder of lakehouse
df = spark.read.format("csv").option("header", "true").schema(orde

# Display the first 10 rows of the dataframe to preview your data
display(df.head(10))
```

14 sec -Apache Spark session ready in 9 sec 788 ms. Command executed in 3 sec 833 PySpark (Python)

The output:

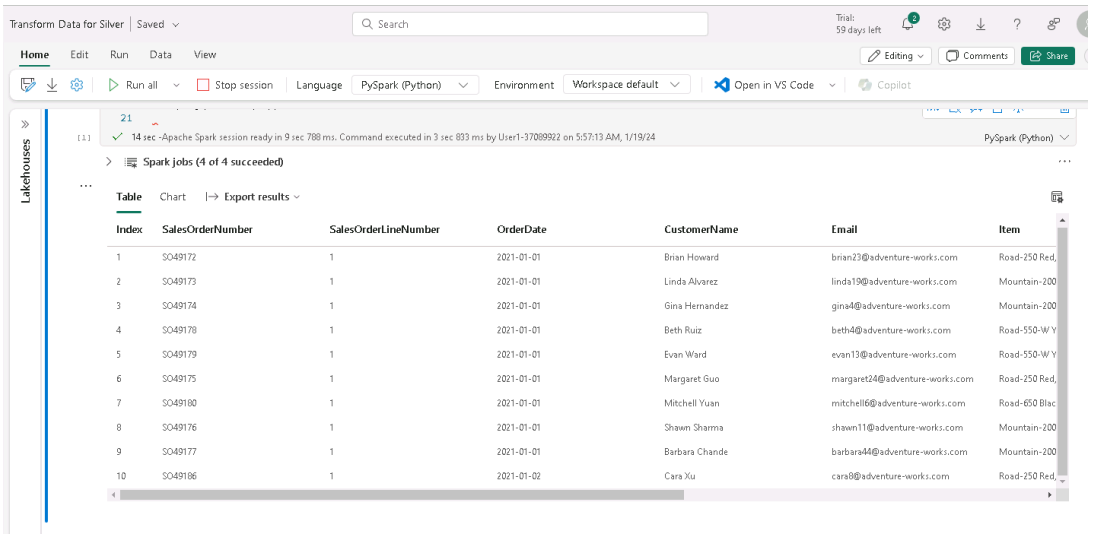

b. Now adding columns for data validation and cleanup. The first line of the code imports the necessary functions from PySpark. then adding new columns to the dataframe to track the source file name, whether the order was flagged as

being a before the fiscal year of interest, and when the row was created and modified.

```
from pyspark.sql.functions import when, lit, col, current_timestamp, input_file_name

# Add columns IsFlagged, CreatedTS and ModifiedTS
df = df.withColumn("FileName", input_file_name()) \
    .withColumn("IsFlagged", when(col("OrderDate") < '2019-08-01',True).otherwise(False)) \
    .withColumn("CreatedTS", current_timestamp()).withColumn("ModifiedTS", current_timestamp())

# Update CustomerName to "Unknown" if CustomerName null or empty
df = df.withColumn("CustomerName", when((col("CustomerName").isNull() | (col("CustomerName")==")),lit("Unknown")).otherwise(col("CustomerN
```

<1 sec -Command executed in 300 ms by User1-37089922 on 6:17:02 AM, 1/19/24

c. Now define the scheme for the sales silver table.

```
# Define the schema for the sales_silver table

from pyspark.sql.types import *
from delta.tables import *

DeltaTable.createIfNotExists(spark) \
    .tableName("sales.sales_silver") \
    .addColumn("SalesOrderNumber", StringType()) \
    .addColumn("SalesOrderLineNumber", IntegerType()) \
    .addColumn("OrderDate", DateType()) \
    .addColumn("CustomerName", StringType()) \
    .addColumn("Email", StringType()) \
    .addColumn("Item", StringType()) \
    .addColumn("Quantity", IntegerType()) \
    .addColumn("UnitPrice", FloatType()) \
    .addColumn("Tax", FloatType()) \
    .addColumn("FileName", StringType()) \
    .addColumn("IsFlagged", BooleanType()) \
    .addColumn("CreatedTS", DateType()) \
    .addColumn("ModifiedTS", DateType()) \
    .execute()
```

PySpark (Python)

d. New sales silver table was created. The triangle indicates that it's a delta table.

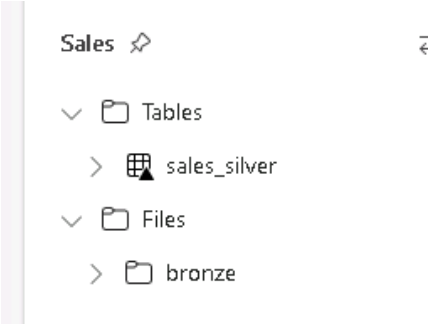

e. perform an upsert operation on a Delta table, updating existing records based on specific conditions and inserting new records when no match is found. Add a new code block. This operation is important because it enables you to update existing

records in the table based on the values of specific columns, and insert new records when no match is found.

```
# Update existing records

from delta.tables import *

deltaTable = DeltaTable.forPath(spark, 'Tables/sale

dfUpdates = df

deltaTable.alias('silver') \
  .merge(
    dfUpdates.alias('updates'),
    'silver.SalesOrderNumber = updates.SalesOrderNu
  ) \
   .whenMatchedUpdate(set =
    {

    }
  ) \
 .whenNotMatchedInsert(values =
    {
      "SalesOrderNumber": "updates.SalesOrderNumber
      "SalesOrderLineNumber": "updates.SalesOrderLi
      "OrderDate": "updates.OrderDate",
      "CustomerName": "updates.CustomerName",
      "Email": "updates.Email",
      "Item": "updates.Item",
      "Quantity": "updates.Quantity",
      "UnitPrice": "updates.UnitPrice",
      "Tax": "updates.Tax",
      "FileName": "updates.FileName",
      "IsFlagged": "updates.IsFlagged",
      "CreatedTS": "updates.CreatedTS",
      "ModifiedTS": "updates.ModifiedTS"
    }
  ) \
   .execute()
```

4. Explore data in the silver layer using the SQL endpoint

   a. This query calculates the total sales for each year in the sales_silver table.

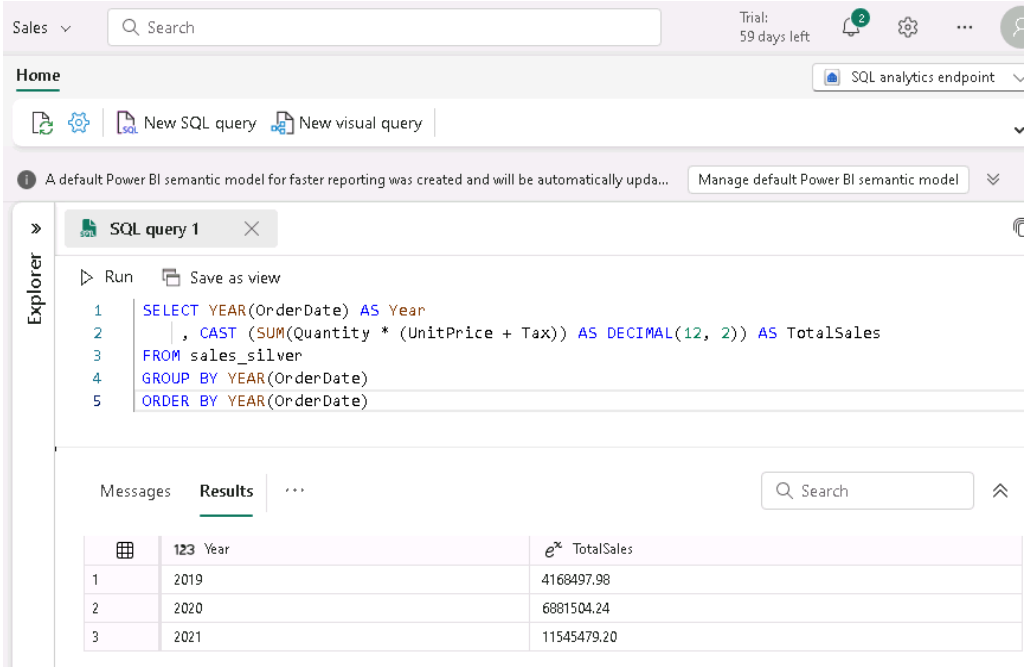

b. This query calculates the total quantity of items purchased by each customer in the sales_silver table, and then returns the top 10 customers in terms of quantity.

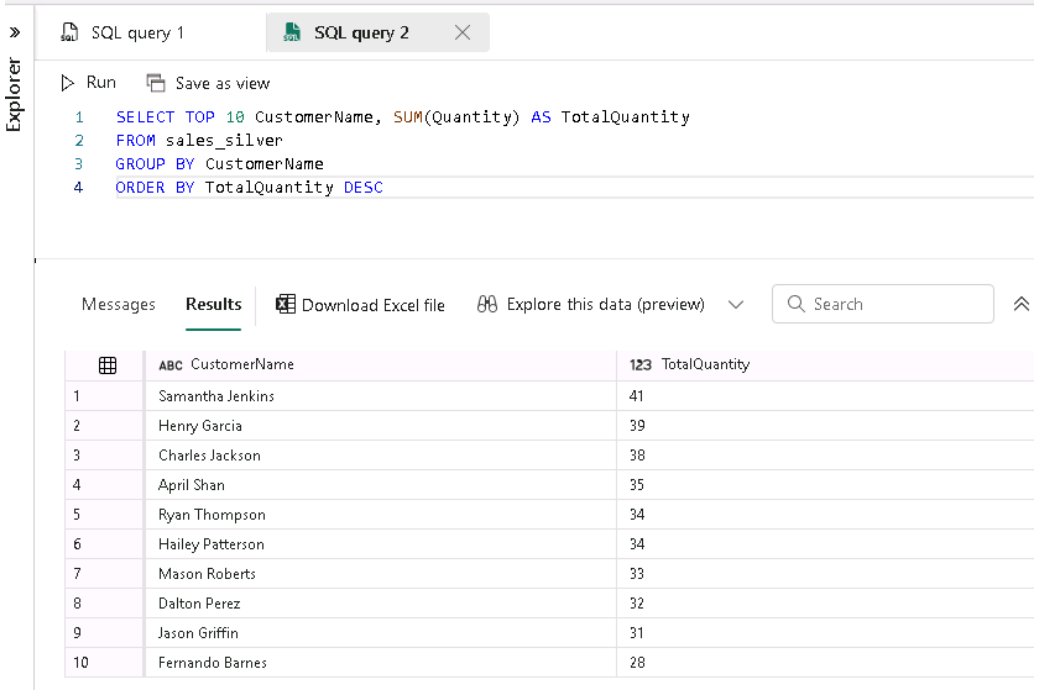

SQL query 1 | SQL query 2

Run | Save as view

```
SELECT TOP 10 CustomerName, SUM(Quantity) AS TotalQuantity
FROM sales_silver
GROUP BY CustomerName
ORDER BY TotalQuantity DESC
```

Messages | Results | Download Excel file | Explore this data (preview) | Search

	CustomerName	TotalQuantity
1	Samantha Jenkins	41
2	Henry Garcia	39
3	Charles Jackson	38
4	April Shan	35
5	Ryan Thompson	34
6	Hailey Patterson	34
7	Mason Roberts	33
8	Dalton Perez	32
9	Jason Griffin	31
10	Fernando Barnes	28

4. Transform Data for Gold Layer

a. Create a new notebook named transform data for gold and add Sales Lakehouse, this will add sales_silver table under tables.
b. Load schema and then add a new code to define the schema

```
from pyspark.sql.types import *
from delta.tables import*

# Define the schema for the dimdate_gold table
DeltaTable.createIfNotExists(spark) \
    .tableName("sales.dimdate_gold") \
    .addColumn("OrderDate", DateType()) \
    .addColumn("Day", IntegerType()) \
    .addColumn("Month", IntegerType()) \
    .addColumn("Year", IntegerType()) \
    .addColumn("mmmyyyy", StringType()) \
    .addColumn("yyyymm", StringType()) \
    .execute()
```

[2] ✓ 5 sec -Command executed in 4 sec 924 ms by User1-37089922 on 7:15:26 AM, 1/19/24 PySpark (Python)

Spark jobs (4 of 4 succeeded) Log

```
<delta.tables.DeltaTable at 0x7f9cac2694e0>
```



```
# Load data to the dataframe as a starting point to create the gold layer
df = spark.read.table("Sales.sales_silver")
```

[1] ✓ 19 sec -Apache Spark session ready in 9 sec 340 ms. Command executed in 10 sec 94 ms by User1-37089 PySpark (Python)

Spark jobs (1 of 1 succeeded)

c. Created 4 tables which was curated, modelled gold layer that can be used for reporting and analysis.

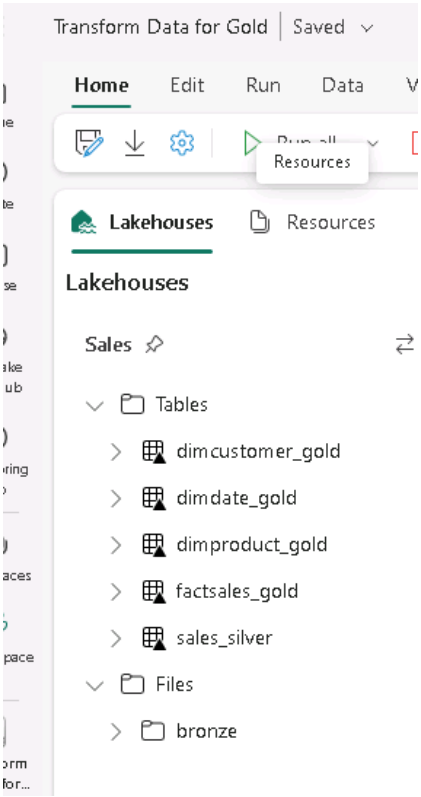

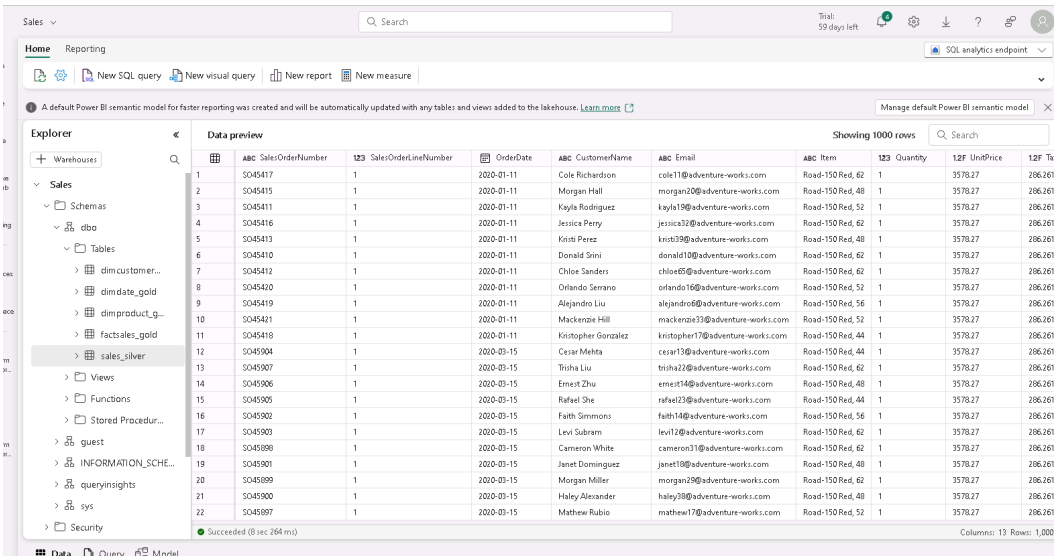

5. Create a database

   a. create a new dataset that includes the gold tables.

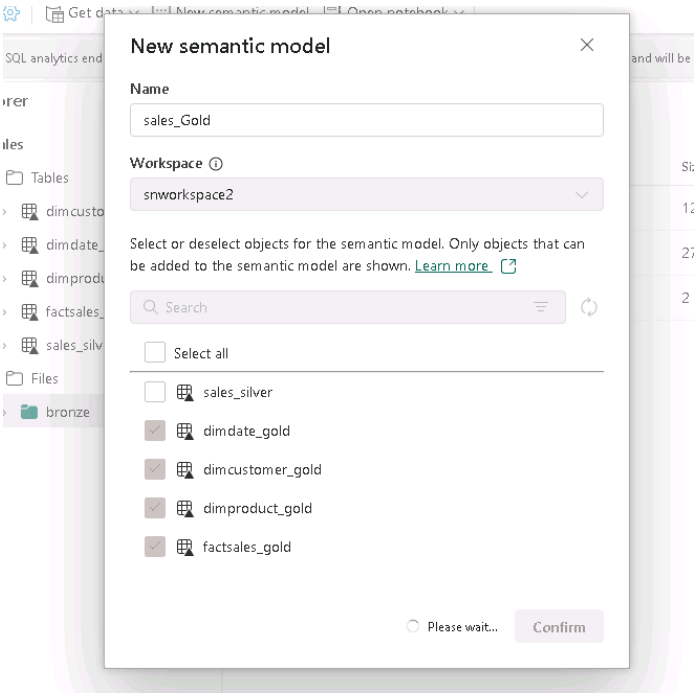

B. set relationship.

The common relationship is sharing order date and it is many to 1.

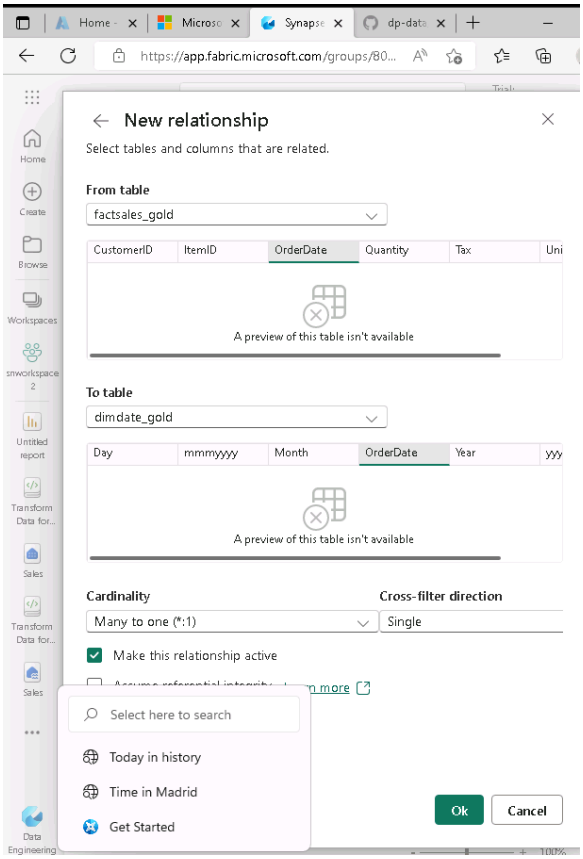

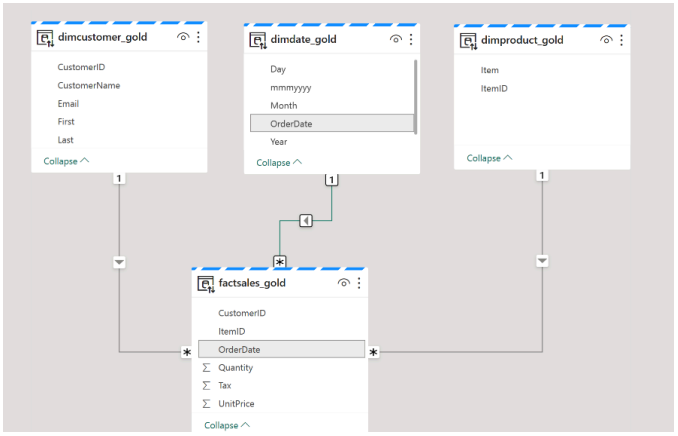
dimcustomer_gold
CustomerID
CustomerName
Email
First
Last
Collapse
dimdate_gold
Day
mmmyyyy
Month
OrderDate
Year
Collapse
dimproduct_gold
Item
ItemID
Collapse
factsales_gold
CustomerID
ItemID
OrderDate
Quantity
Tax
UnitPrice
Collapse
1
1
1
*
*
*